

Chapter 2

Application Layer

A note on the use of these Powerpoint slides:

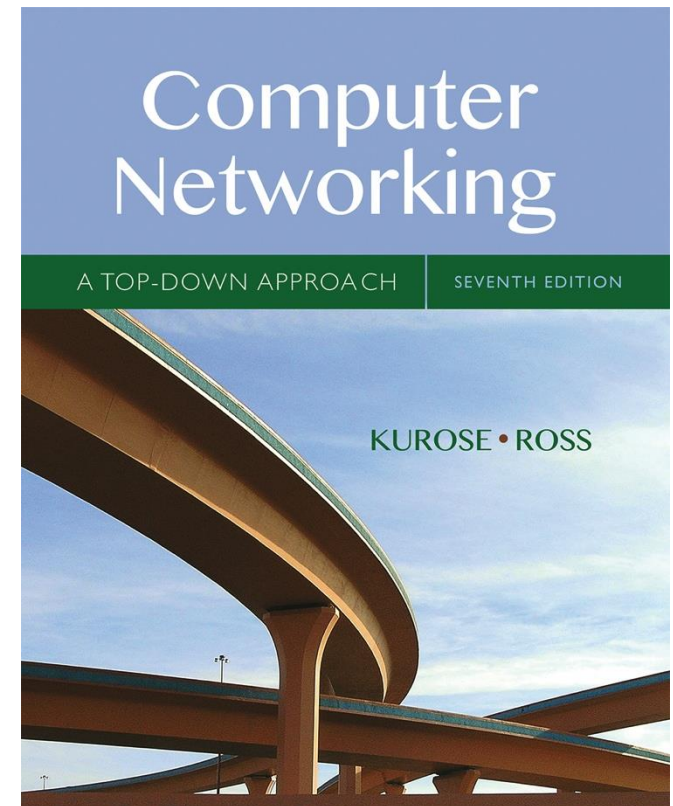
We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016

J.F Kurose and K.W. Ross, All Rights Reserved



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

Some network apps

- e-mail
- web
- text messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube, Hulu, Netflix)
- voice over IP (e.g., Skype)
- real-time video conferencing
- social networking
- search
- ...
- ...

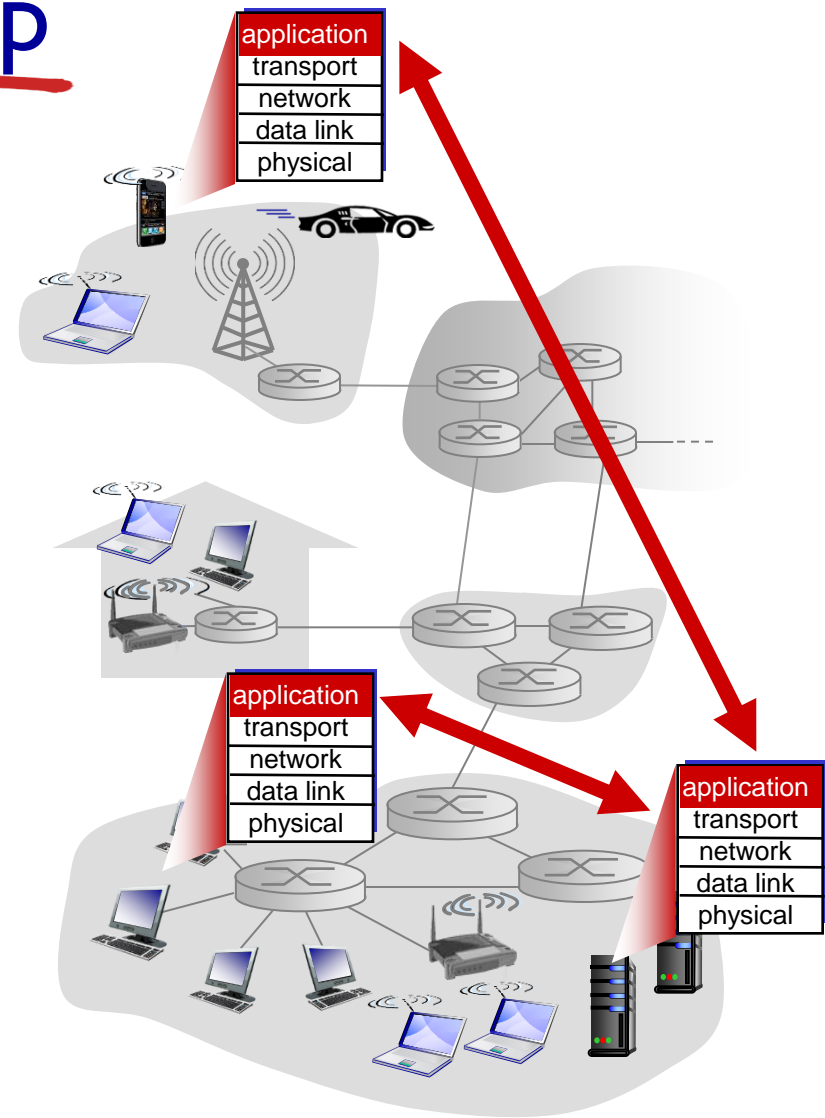
Creating a network app

write programs that:

- run on (different) *end systems*
- communicate over network
- e.g., web server software communicates with browser software

no need to write software
for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation




Application architectures

possible structure of applications:

- client-server
- peer-to-peer (P2P)

Client-server architecture



server:


- always-on host
- permanent IP address
- data centers for scaling

clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other

P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
 - complex management



Processes communicating

process: program running within a host

- within same host, two processes communicate using **inter-process communication** (defined by OS)
- processes in different hosts communicate by exchanging **messages**

clients, servers

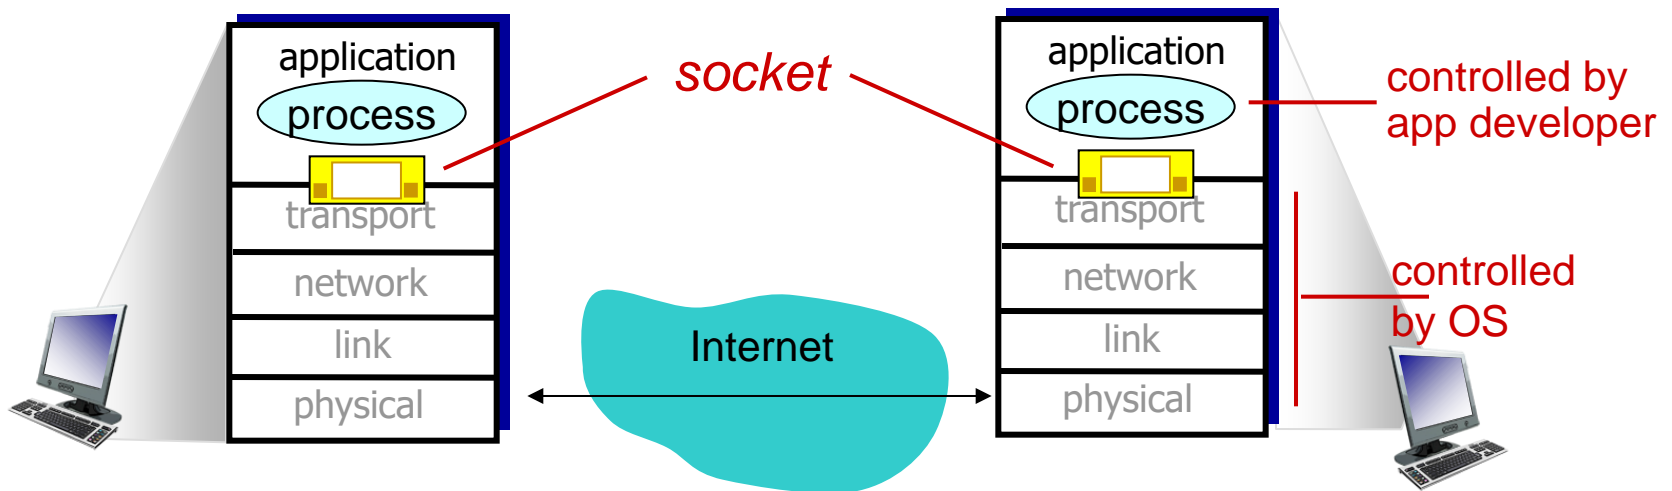
client process: process that initiates communication

server process: process that waits to be contacted

- aside: applications with P2P architectures have client processes & server processes

Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
 - sending process shoves message out door
 - sending process relies on **transport infrastructure** on other side of door to deliver message to socket at receiving process



Addressing processes

- to receive messages, process must have *identifier*
- host device has unique 32-bit IP address
- Q: does IP address of host on which process runs suffice for identifying the process?
 - A: no, *many* processes can be running on same host
- *identifier* includes both **IP address** and **port numbers** associated with process on host.
- example port numbers:
 - HTTP server: 80
 - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
 - **IP address**: 128.119.245.12
 - **port number**: 80
- more shortly...

App-layer *protocol* defines

- types of messages exchanged,
 - e.g., request, response
- message syntax:
 - what fields in messages & how fields are delineated
- message semantics
 - meaning of information in fields
- rules for when and how processes send & respond to messages

open protocols:

- defined in RFCs
- allows for interoperability
- e.g., HTTP, SMTP

proprietary protocols:

- e.g., Skype

What transport service does an app need?

data loss

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- other apps (“elastic apps”) make use of whatever throughput they get

Transport service requirements: common apps

| application | data loss | throughput | time sensitive |
|-----------------------|---------------|---|------------------|
| file transfer | no loss | elastic | no |
| e-mail | no loss | elastic | no |
| Web documents | no loss | elastic | no |
| real-time audio/video | loss-tolerant | audio: 5kbps-1Mbps video: 10kbps-5Mbps | yes, 100' s msec |
| stored audio/video | loss-tolerant | same as above | yes, few secs |
| interactive games | loss-tolerant | few kbps up | yes, 100' s msec |
| smartphone messaging | no loss | elastic | yes and no |

Internet transport protocols services

TCP service:

- *reliable transport* between sending and receiving process (in order delivery, recovery of lost data)
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum throughput guarantee, security
- *connection-oriented*: setup required between client and server processes

UDP service:

- *unreliable data transfer* between sending and receiving process
- *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

Q: why bother? Why is there a UDP?

*More about transport
protocols in chapter#3*

Internet apps: application, transport protocols

| application | application layer protocol | underlying transport protocol |
|------------------------|---|--------------------------------------|
| e-mail | SMTP [RFC 2821] | TCP |
| remote terminal access | Telnet [RFC 854] | TCP |
| Web | HTTP [RFC 2616] | TCP |
| file transfer | FTP [RFC 959] | TCP |
| streaming multimedia | HTTP (e.g., YouTube), RTP [RFC 1889] | TCP or UDP |
| Internet telephony | SIP, RTP, proprietary (e.g., Skype) | TCP or UDP |

More about transport layer protocols in chapter 3....

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

Web and HTTP

First, a review...

- *web page* consists of *objects*
- object can be HTML file, JPEG image, audio file,...
- web page consists of *base HTML-file* which includes *several referenced objects*
- each object is addressable by a *URL*, e.g.,

`www.someschool.edu/someDept/pic.gif`

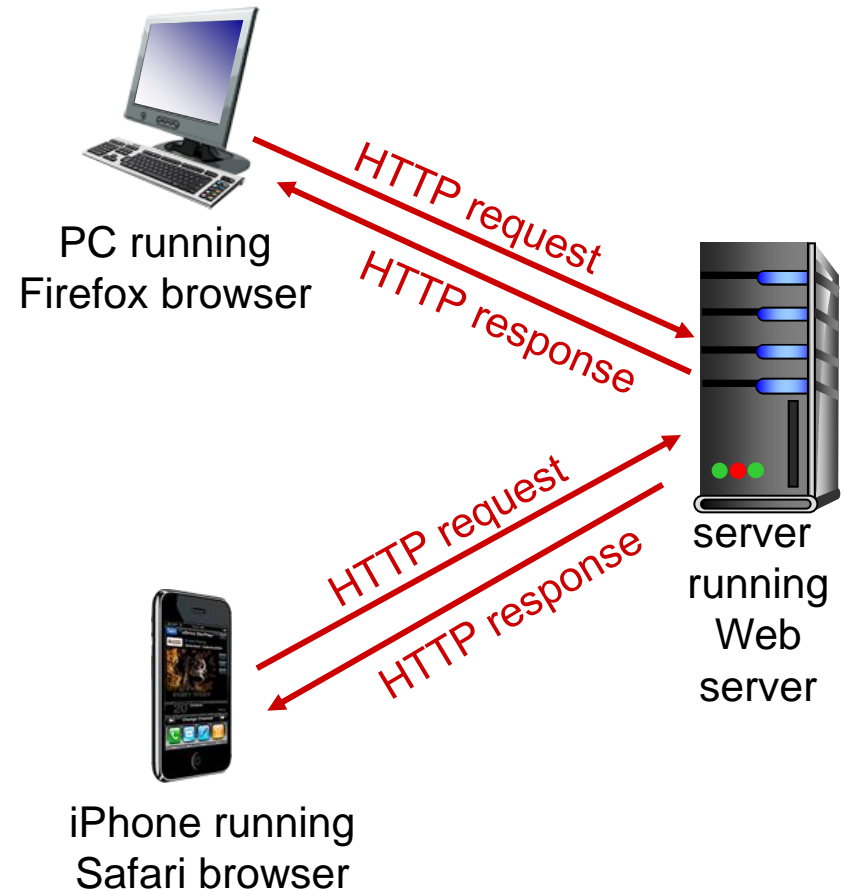
host name

path name

HTTP overview

HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
 - **client**: browser that requests, receives, (using HTTP protocol) and "displays" Web objects
 - **server**: Web server sends (using HTTP protocol) objects in response to requests



HTTP overview (continued)

uses TCP:

1. client initiates TCP connection (creates socket) to server, port 80
2. server accepts TCP connection from client
3. HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
4. TCP connection closed

HTTP connections

non-persistent HTTP

- at most one object sent over TCP connection
 - connection then closed
- downloading multiple objects required multiple connections

persistent HTTP

- multiple objects can be sent over single TCP connection between client, server, serially.
- reduces TCP overhead

Non-persistent HTTP (Serial)

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,
references to 10
jpeg images)

1a. HTTP client initiates **TCP connection** to HTTP server (process) at `www.someSchool.edu` on port 80

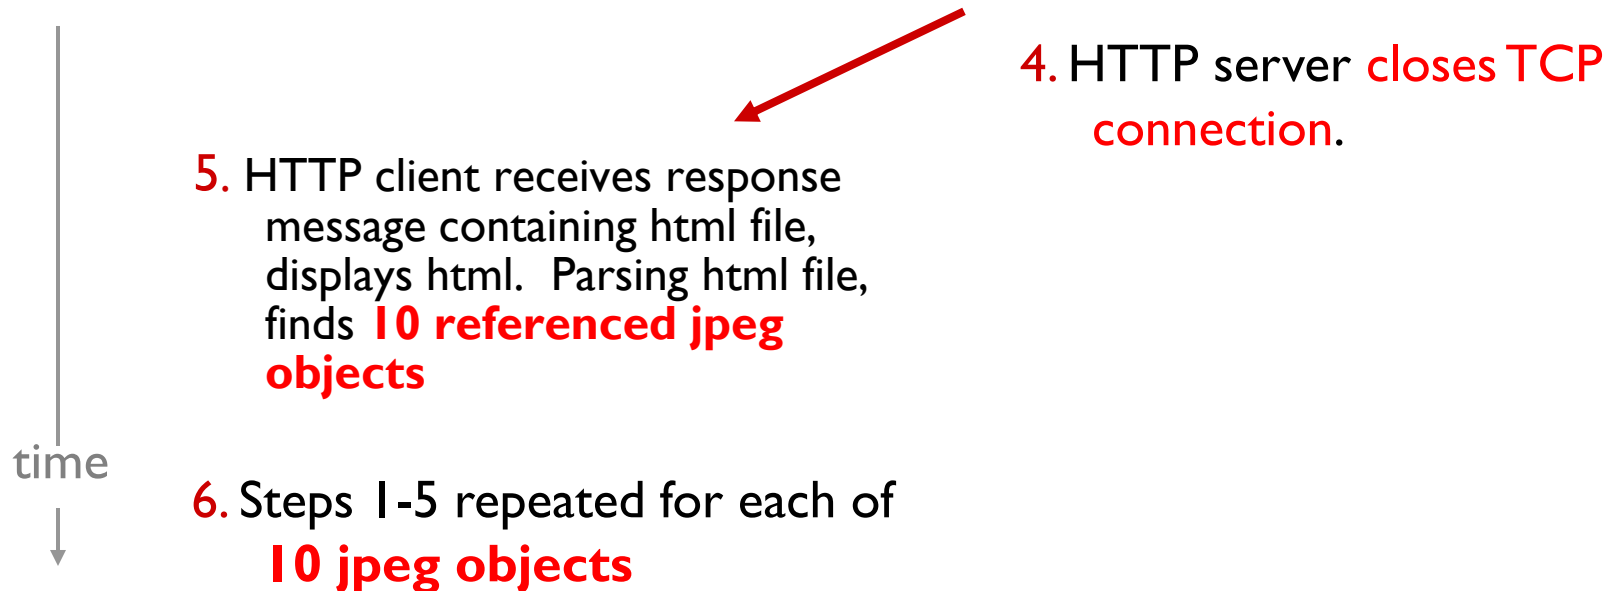
1b. HTTP server at host `www.someSchool.edu` waiting for **TCP connection** at port 80. **“accepts”** connection, notifying client

2. HTTP client sends HTTP **request message** (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms **response message** containing requested object, and sends message into its socket

time
↓

Non-persistent HTTP (cont.)



Non-persistent HTTP: response time

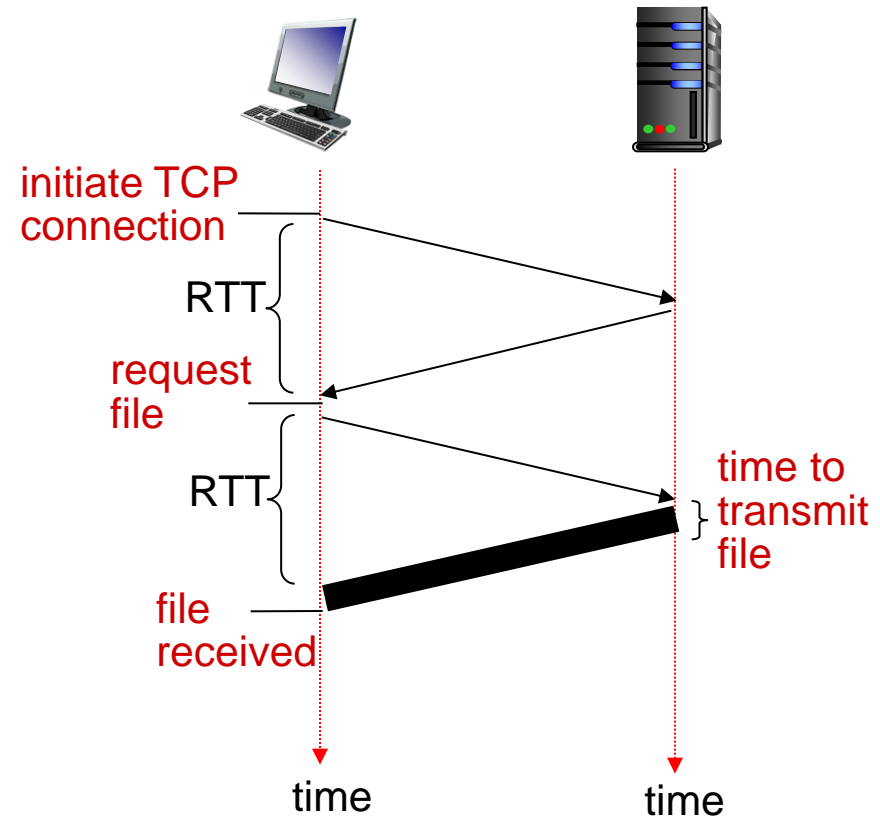
RTT (definition): time for a *small packet* to travel from client to server and back

HTTP response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =

$2\text{RTT} + \text{file transmission time}$

File transmission time = ?



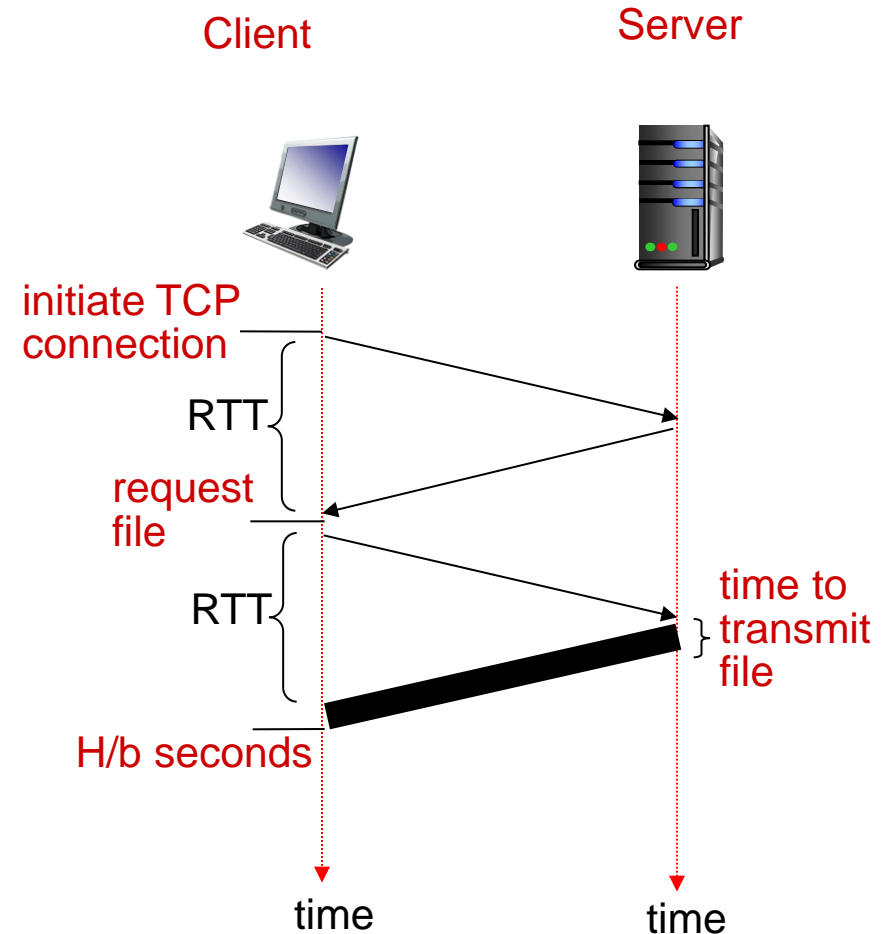
Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =

$2\text{RTT} + \text{file transmission time}$



Persistent HTTP

non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for each TCP connection

persistent HTTP:

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- For example, it will take total of 6 RTTs to download a webpage containing 4 objects.

Class Exercise

- *Consider a client fetching 5 objects from a server. Assume that all objects are of same sizes (B bits) and the client is using persistent HTTP connection. The transmission rate is R bits/second. Draw the timeline diagram for the request response events between client and server and find out the total delay.*

HTTP request message

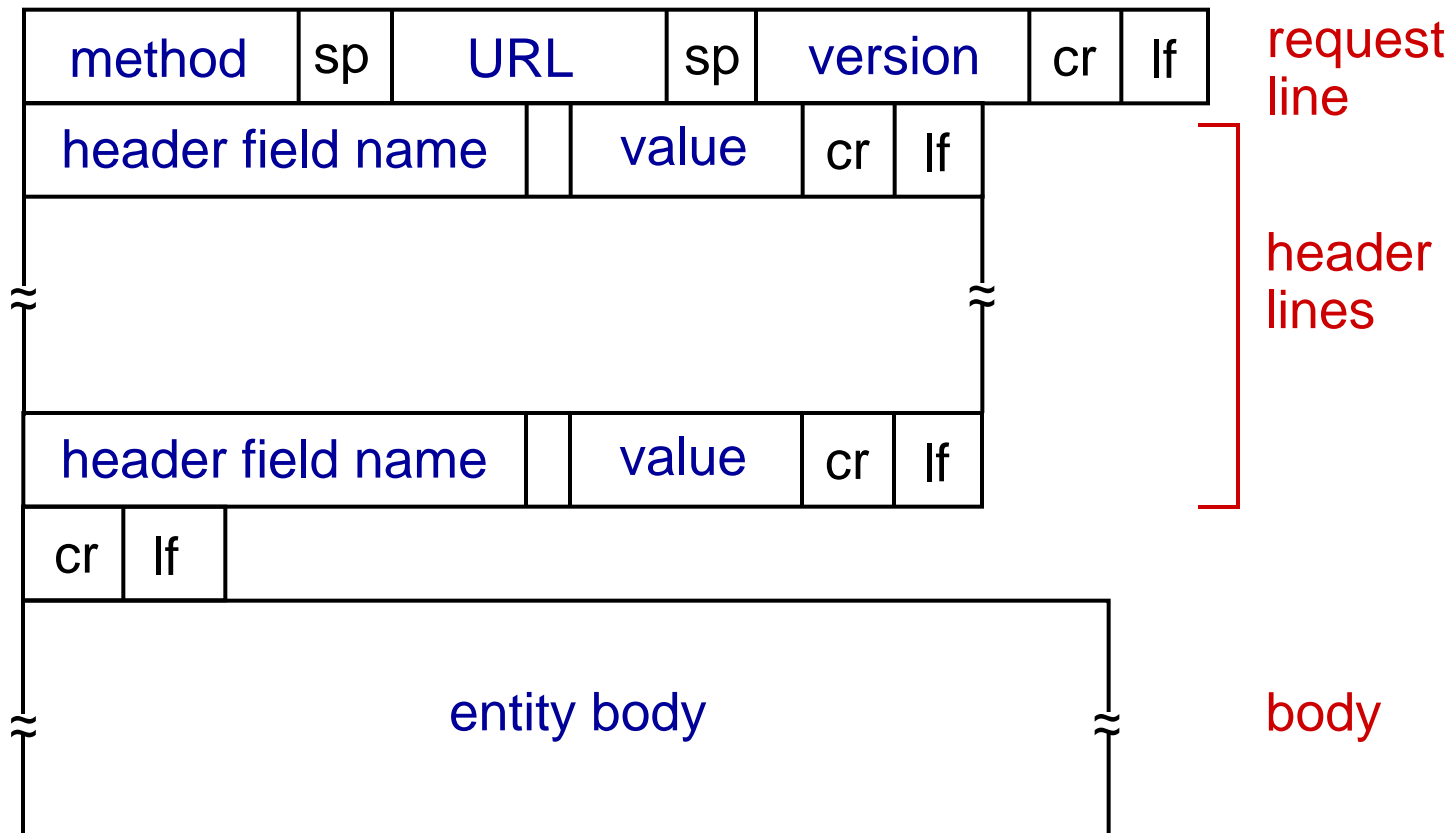
- two types of HTTP messages: *request, response*
- **HTTP request message:**

The diagram illustrates the structure of an HTTP request message. It consists of a request line followed by header lines, each terminated by a carriage return and line-feed character sequence (\r\n). Annotations with arrows point to specific parts of the message:

- request line (GET, POST, HEAD commands)**: Points to the first line of the message.
- header lines**: Points to the block of lines containing the request headers.
- carriage return, line feed at start of line indicates end of header lines**: Points to the \r\n sequence at the end of the header block.
- carriage return character**: Points to the \r character in the \r\n sequence at the end of the request line.
- line-feed character**: Points to the \n character in the \r\n sequence at the end of the request line.

```
GET /index.html HTTP/1.1\r\nHost: www.someschool.edu\r\nUser-Agent: Firefox/3.6.10\r\nAccept-Language: en-us,en;q=0.5\r\nConnection: close\r\n\r\n
```

HTTP request message: general format



Method types

- **GET:** retrieve a file
- **POST:** submitting a form to a server
- **HEAD:** asks server to leave requested object out of response, and just get meta-data
- **Put:** uploads file in entity body to path specified in URL
- **Delete:** remove named source

HTTP response message

status line
(protocol
status code
status phrase)

header
lines

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
      GMT\r\n
Content-Length: 2652\r\n
Connection: close\r\n
Content-Type: text/html\r\n
\r\n
```

data, e.g.,
requested
HTML file

data data data data data ...

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this msg

301 Moved Permanently

- requested object moved, new location specified later in this msg (Location:)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

User-server state: cookies

many Web sites use cookies

example:

- Susan always access Internet from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
 - unique ID
 - entry in backend database for ID

Cookies: keeping “state” (cont.)

client



server



cookie file



ebay 8734
amazon 1678


usual http request msg

Amazon server
creates ID
1678 for user

usual http response
set-cookie: 1678

create
entry

backend
database



usual http request msg
cookie: 1678

cookie-
specific
action

access

usual http response msg

access

cookie-
specific
action

one week later:



ebay 8734
amazon 1678

usual http request msg
cookie: 1678

usual http response msg

User-server state: cookies

many Web sites use cookies

four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

Cookies (continued)

what cookies can be used for:

- authorization
- shopping carts
- recommendations

aside

cookies and privacy:

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites


how to keep “state”:

- **protocol endpoints:** maintain state at sender/receiver over multiple transactions
- **cookies:** http messages carry state

Web caches (proxy server)

goal: satisfy client request without involving origin server

- user sets browser: Web accesses via cache
- browser sends all HTTP requests to cache
 - **If object in cache:** cache returns object
 - else cache requests object from origin server, then returns object to client



More about Web caching

- cache acts as both client and server
 - server for original requesting client
 - client to origin server
- typically cache is installed by ISP (university, company, residential ISP)

why Web caching?

- reduce response time for client request
- reduce traffic on an institution's access link
- Internet dense with caches: enables “poor” content providers to effectively deliver content

Caching example:


assumptions:

- avg object size: 1Mbits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 100 Mbps
- RTT from *internet router* to any origin server: 2 sec
- access link rate: 15 Mbps

consequences:

- LAN utilization: 1%
- access link utilization = 100%
- total delay = Internet delay + access delay + LAN delay
= 2 sec + seconds + msecs

≥ 3 seconds




Caching example: fatter access link

assumptions:

- avg object size: 1Mbits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 100 Mbps
- RTT from internet router to any origin server: 2 sec
- access link rate: ~~15 Mbps~~ → 100 Mbps

consequences:

- LAN utilization: 1%
- access link utilization = ~~100%~~ → 15%
- total delay = Internet delay + access delay + LAN delay
= 2 sec + ~~seconds~~ → msecs



Cost: increased access link speed (not cheap!)

Caching example: install local cache


assumptions:

- avg object size: 1M bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 100 Mbps
- RTT from internet router to any origin server: 2 sec
- access link rate: 15 Mbps

consequences:

- LAN utilization: 1%
- access link utilization = **100%**
- total delay = Internet delay + access delay + LAN delay
= 0 + 0 + 15 msecs


Cost: web cache (cheap!)



Caching example: install local cache

Calculating access link utilization, delay with cache:

- suppose cache hit rate is 0.4
 - 40% requests satisfied at cache, 60% requests satisfied at origin
- total delay
 - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
 - $= 0.6 (3.15) + 0.4 (0.15) = \sim 1.95 \text{ secs}$
 - less than with 100 Mbps **access link** (and cheaper too!)



Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version
 - no object transmission delay
 - lower link utilization
- **cache:** specify date of cached copy in HTTP request

If-modified-since: <date>
- **server:** response contains no object if cached copy is up-to-date:

HTTP/1.0 304 Not Modified

(Proxy Server)
client



Origin server

