

Chapter 2

Application Layer

A note on the use of these Powerpoint slides:

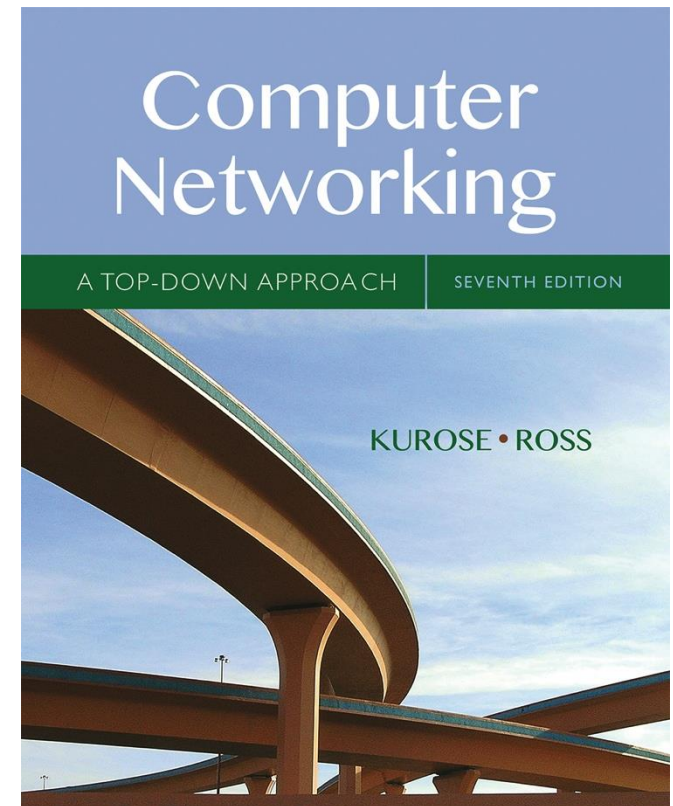
We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016

J.F Kurose and K.W. Ross, All Rights Reserved



Computer Networking: A Top Down Approach

7th edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7

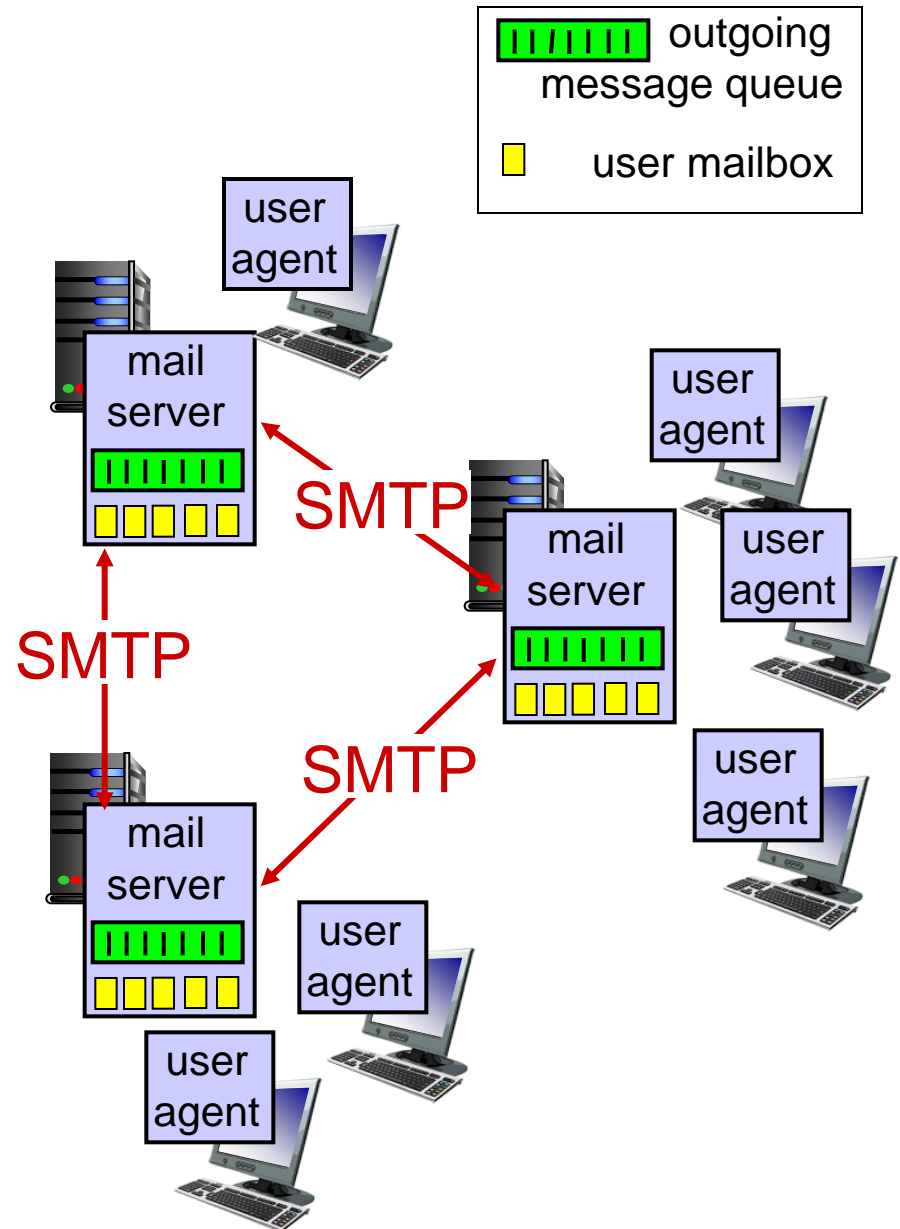
Electronic mail

Three major components:

- user agents
- mail servers
- simple mail transfer protocol: SMTP

User Agent

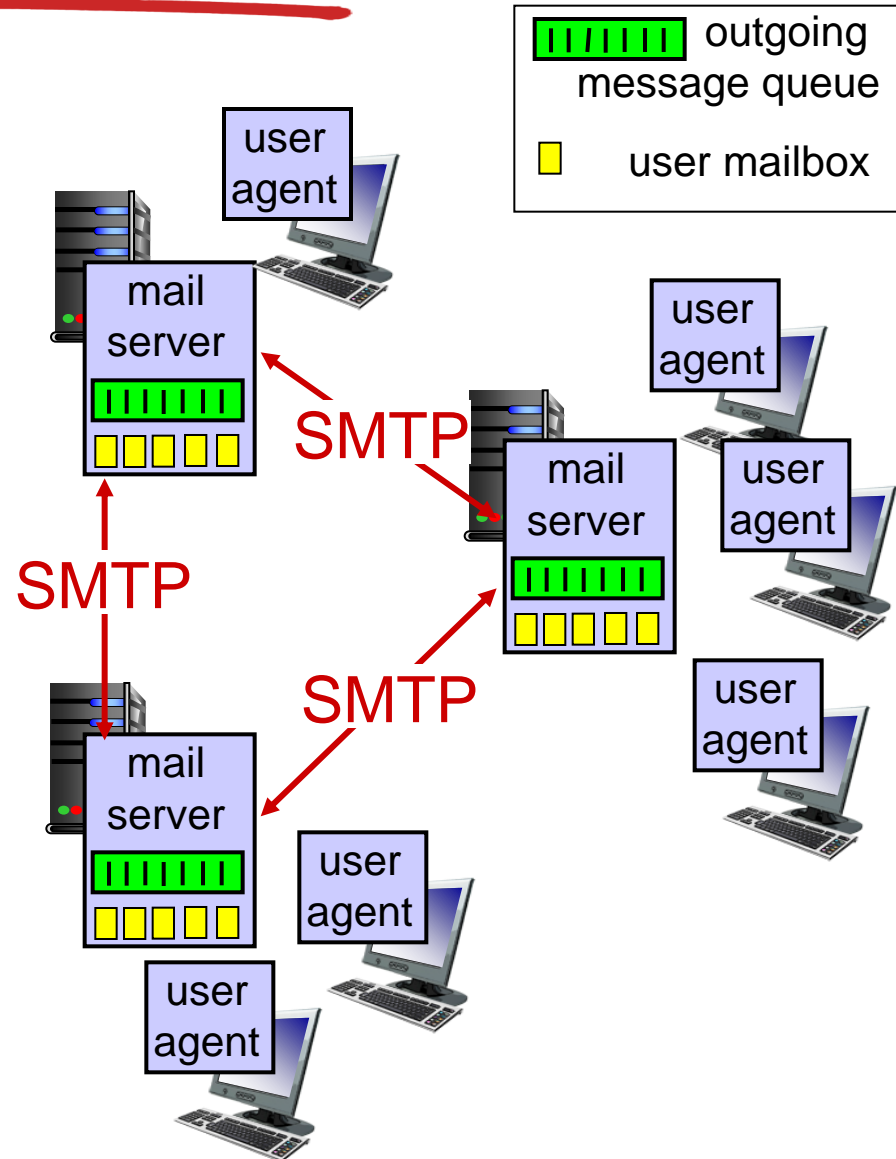
- a.k.a. “mail reader”
- composing, editing, reading mail messages
- e.g., Outlook, Thunderbird, iPhone mail client
- outgoing, incoming messages stored on server



Electronic mail: mail servers

mail servers:

- *mailbox* contains incoming messages for user
- *message queue* of outgoing (to be sent) mail messages
- *SMTP protocol* between mail servers to send email messages
 - **client:** sending mail server
 - **server:** receiving mail server

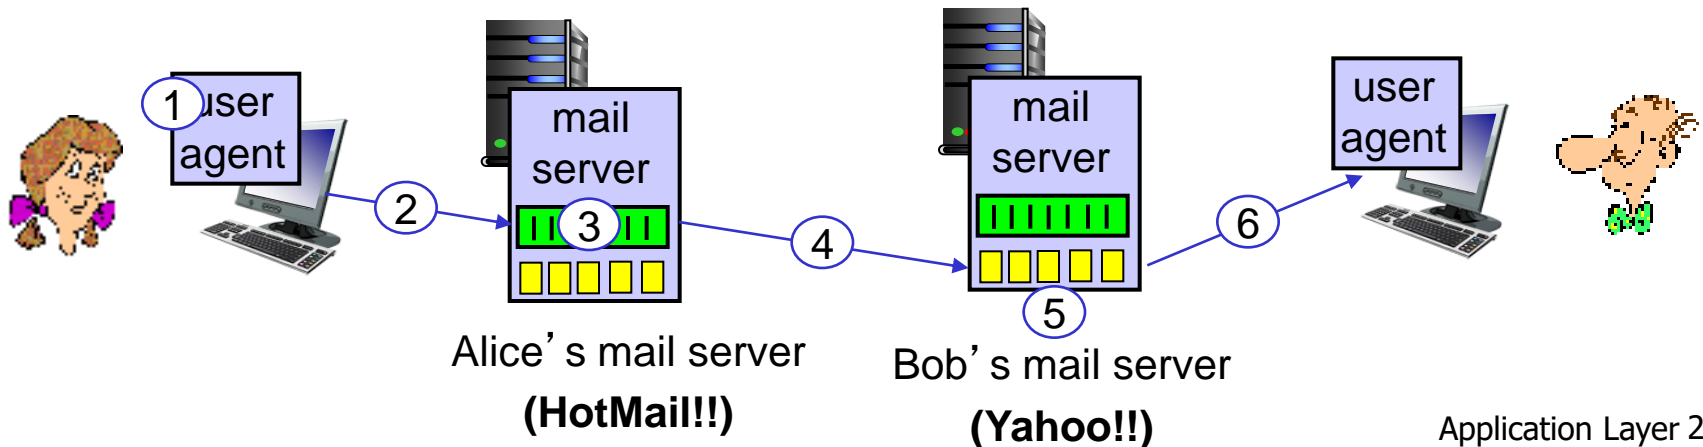


Electronic Mail: SMTP [RFC 2821]

- uses TCP to reliably transfer email message from client to server, port 25
- **direct transfer**: sending server to receiving server
- three phases of transfer
 - handshaking (greeting)
 - transfer of messages
 - closure
- command/response interaction (like HTTP)
 - **response**: status code and phrase

Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message “to”
`Bob@yahoo.com`
- 2) Alice’s UA sends message to her mail server; message placed in message queue
- 3) client side of SMTP opens TCP connection with Bob’s mail server
- 4) SMTP client sends Alice’s message over the TCP connection
- 5) Bob’s mail server places the message in Bob’s mailbox
- 6) Bob invokes his user agent to read message



Sample SMTP interaction

- After the TCP connection is established with the server, it sends a 220

```
S: 220 yahoomail.com
C: HELO Hotmail.com
S: 250 Hello Hotmail.com, pleased to meet you
C: MAIL FROM: <Alice@Hotmail.com>
S: 250 Alice@Hotmail.com... Sender ok
C: RCPT TO: <Bob@yahoomail.com>
S: 250 Bob@yahoomail.com ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Hi, How are you?
C: This is just a SMTP testing.
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 yahoomail.com closing connection
```

SMTP: final words

- SMTP uses persistent connections
 - All the emails are sent over a single TCP connection.

comparison with HTTP:


- HTTP: pull
- SMTP: push
 - the TCP connection is initiated by the machine that wants to send the file
- **HTTP:** each object encapsulated in its own response message
- **SMTP:** places all of the message's objects into one message.

Mail message format

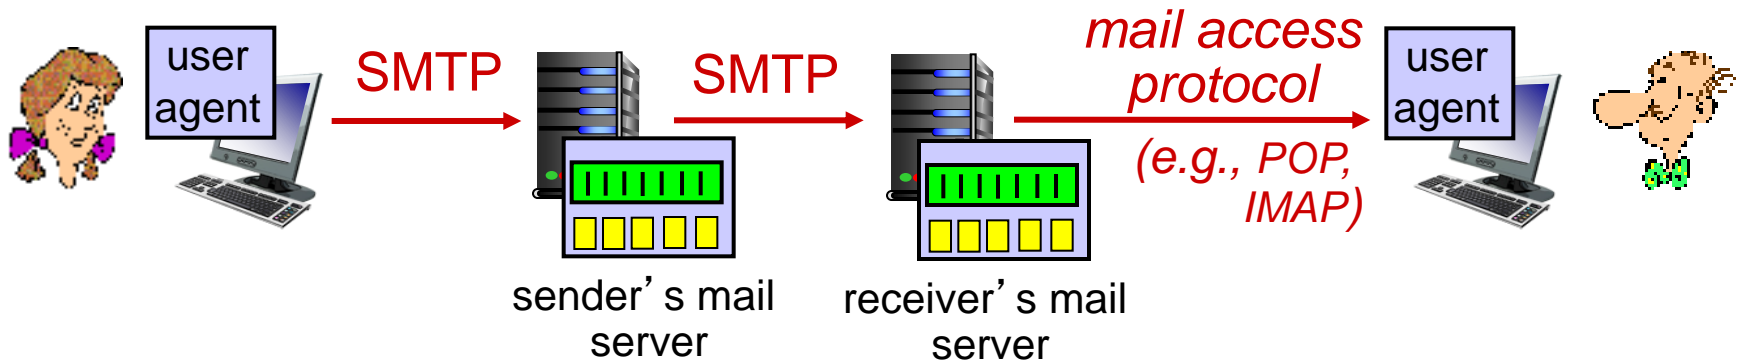
SMTP: protocol for exchanging email messages

RFC 822: standard for text message format:

- header lines, e.g.,
 - To:
 - From:
 - Subject:
- Body: the “message”
 - ASCII characters only



Mail access protocols



- **SMTP:** delivery/storage to receiver's server
- **mail access protocol:** retrieval from server
 - **POP:** Post Office Protocol [RFC 1939]: authorization, download
 - **IMAP:** Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored messages on server

POP3 protocol

authorization phase

- client commands:
 - **user**: declare username
 - **pass**: password
- server responses
 - +OK
 - -ERR

transaction phase, client:

- **list**: list message numbers
- **retr**: retrieve message by number
- **dele**: delete
- **quit**

After processing quit command, the POP3 server enters update phase and removes deleted messages from the mailbox.

```
S: +OK POP3 server ready
C: user alice
S: +OK
C: pass student
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

POP3 (more) and IMAP

more about POP3

- POP3 uses port 110
- A user agent using POP3 can often be configured (by the user) to:
 - "download and delete" or
 - to "download and keep".
- In “download and delete” user cannot re-read e-mail if he changes client

IMAP

- IMAP uses port 993
- keeps all messages in one place: at server
- allows user to organize messages in folders
- keeps user state across sessions:
 - names of folders and mappings between message IDs and folder name

Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

DNS: domain name system

people: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

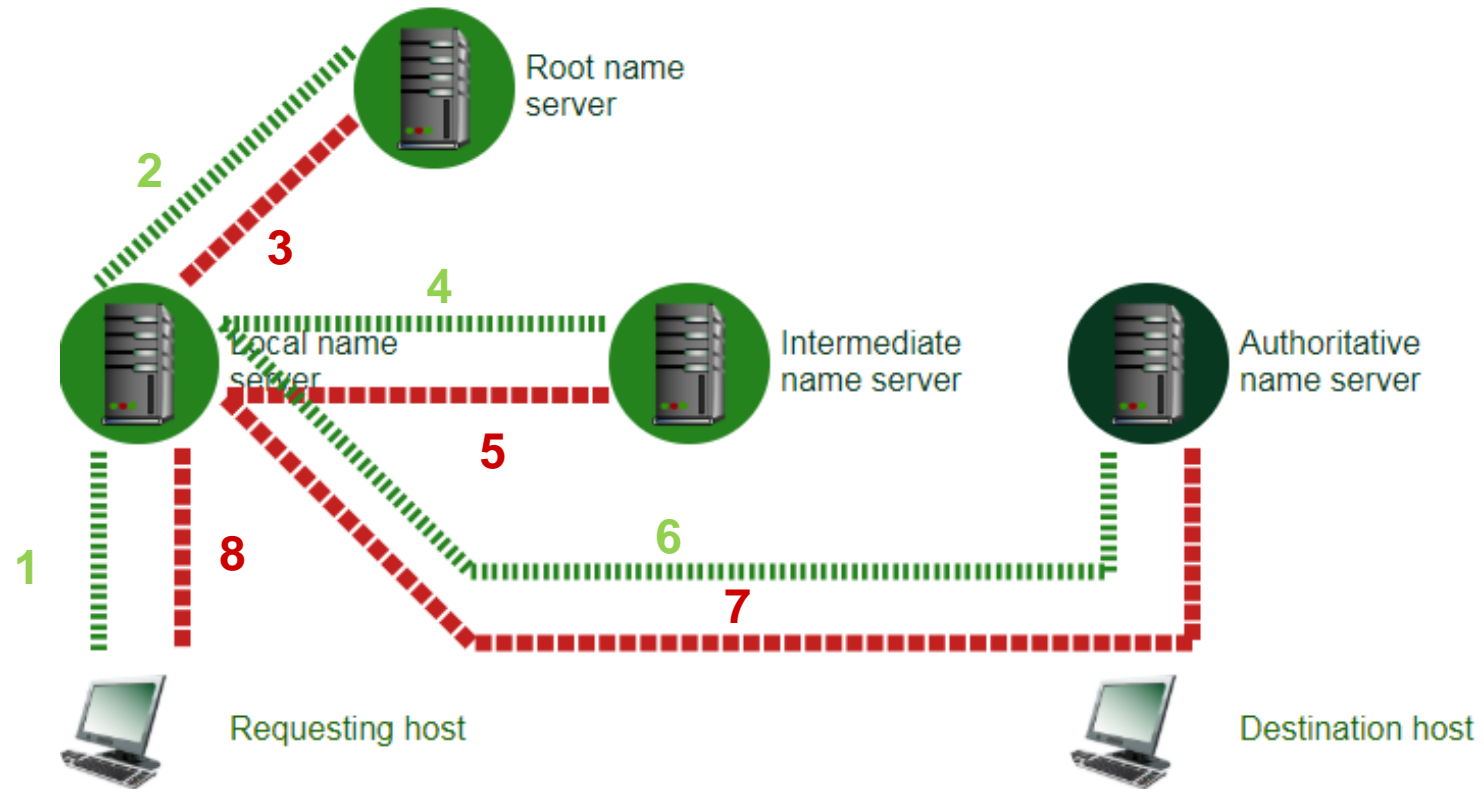
- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., `www.yahoo.com` - used by humans

Q: how to map between IP address and name, and vice versa ?

Domain Name System:

- *distributed database*
implemented in hierarchy of many *name servers*
- *application-layer protocol*: hosts, name servers communicate to *resolve* names (address/name translation)
 - **note**: core Internet function, implemented as application-layer protocol
 - complexity at network's “edge”

Bigger Picture of DNS Operation



DNS: services, structure

DNS services

- hostname to IP address translation
- **host aliasing**
 - Canonical (real), alias names
 - A host with canonical (real) host name (e.g., relay1.west-coast.enterprise.com) may have several alias names (www.enterprise.com, enterprise.com)
- **mail server aliasing, e.g.,**
- Permit both mail/Web server to have identical (aliased) host name
- **someone@yahoo.com** instead of **someone@mail.yahoo.com**
- **load distribution**
 - replicated Web servers: many IP addresses correspond to one name
- **DNS is built over UDP**

why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

A: doesn't scale!

Introduction

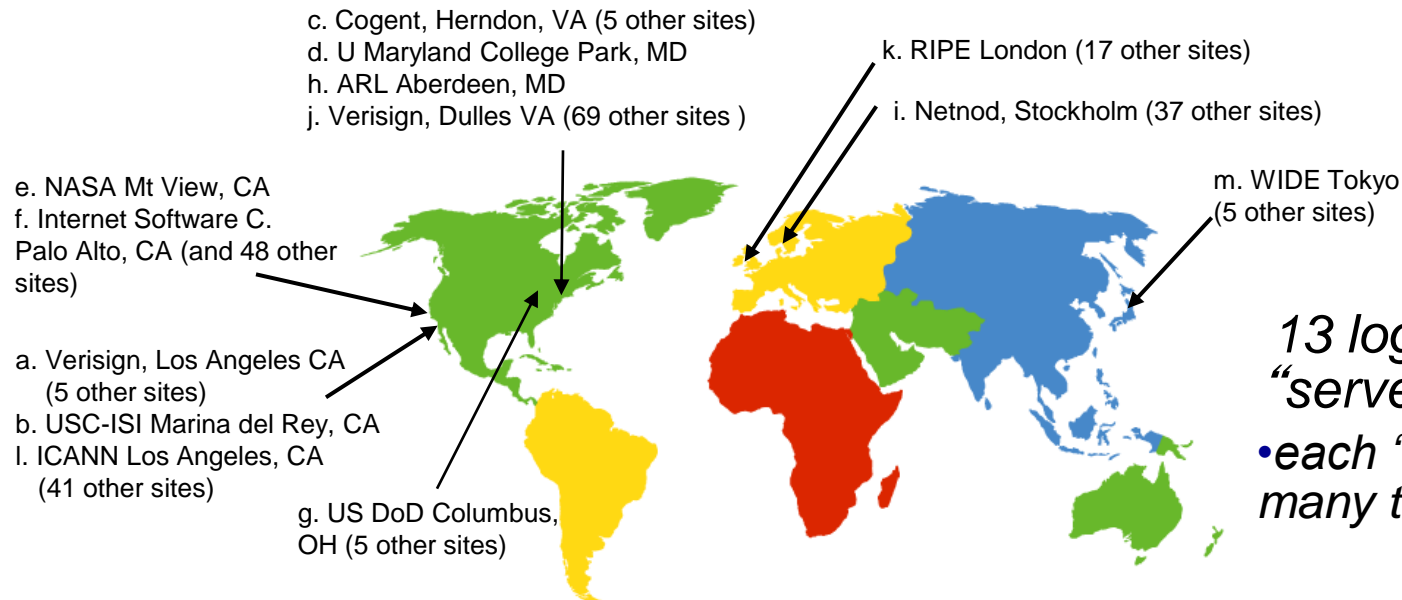
DNS: root name servers

■ Root name servers

- contacted by *local name server* that can not resolve name

■ root name server:

- *13 logical root name “servers” worldwide each “server” replicated many times*
- return IP addresses of TLD servers to local name server



13 logical root name “servers” worldwide
• *each “server” replicated many times*

TLD, authoritative servers

- *top-level domain (TLD) servers (or Intermediate Name Server):*
 - responsible for com, org, net, edu, aero and all top-level country domains, e.g.: uk, fr, ca, jp
 - For example, *Network Solutions* maintains servers for .com TLD
 - and, *Educause* for .edu TLD
- *authoritative DNS servers:*
 - organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
 - can be maintained by organization or service provider

Local DNS name server

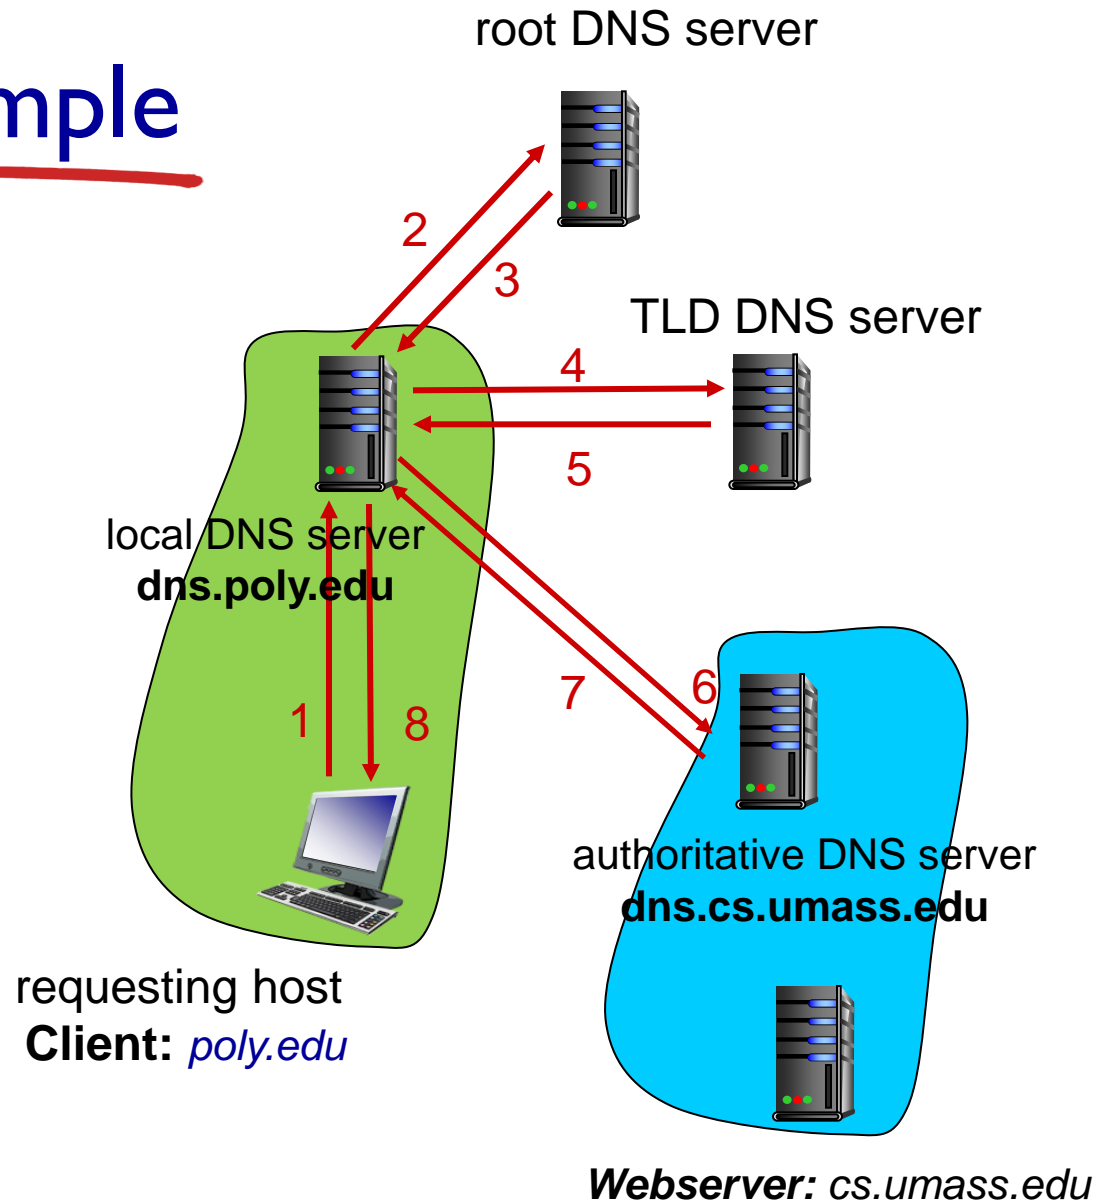
- **Local Name Server**
- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
 - also called “*default name server*”
- when host makes DNS query, query is sent to its local DNS server
 - has *local cache* of recent name-to-address translation pairs (but may be out of date!)
 - acts as proxy, forwards query into hierarchy

DNS name resolution example

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

iterated query:


- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



DNS name resolution example

recursive query:

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy



DNS: caching, updating records

- once (any) name server learns mapping, it *caches* mapping
 - cache entries timeout (disappear) after some time (TTL)
 - TLD servers (intermediate name servers) typically cached in local name servers
 - thus root name servers not often visited
- cached entries may be *out-of-date*
 - if name host changes IP address, may not be known Internet-wide until all TTLs expire

DNS records

DNS: distributed database storing resource records (RR)

RR format: (name, value, type, ttl)

- **Query**(domain name, RR type)
 - Resource Record (RR) type is like an attribute type
- **Answer**(values, additional RRs)

type=A

- **name** is hostname
- **value** is IP address

type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain

type=CNAME

- **name** is alias name for some “canonical” (the real) name
- **www.ibm.com** is really **servereast.backup2.ibm.com**
- **value** is canonical name

type=MX (mail exchanger)

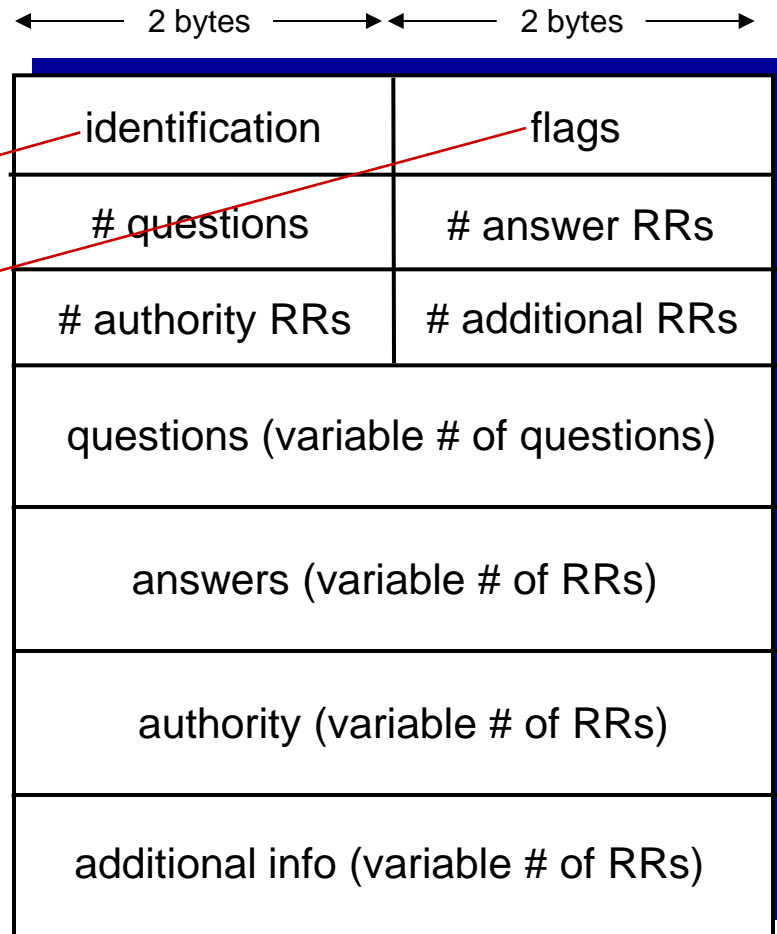
- **value** is name of mailserver associated with **name**

DNS protocol, messages

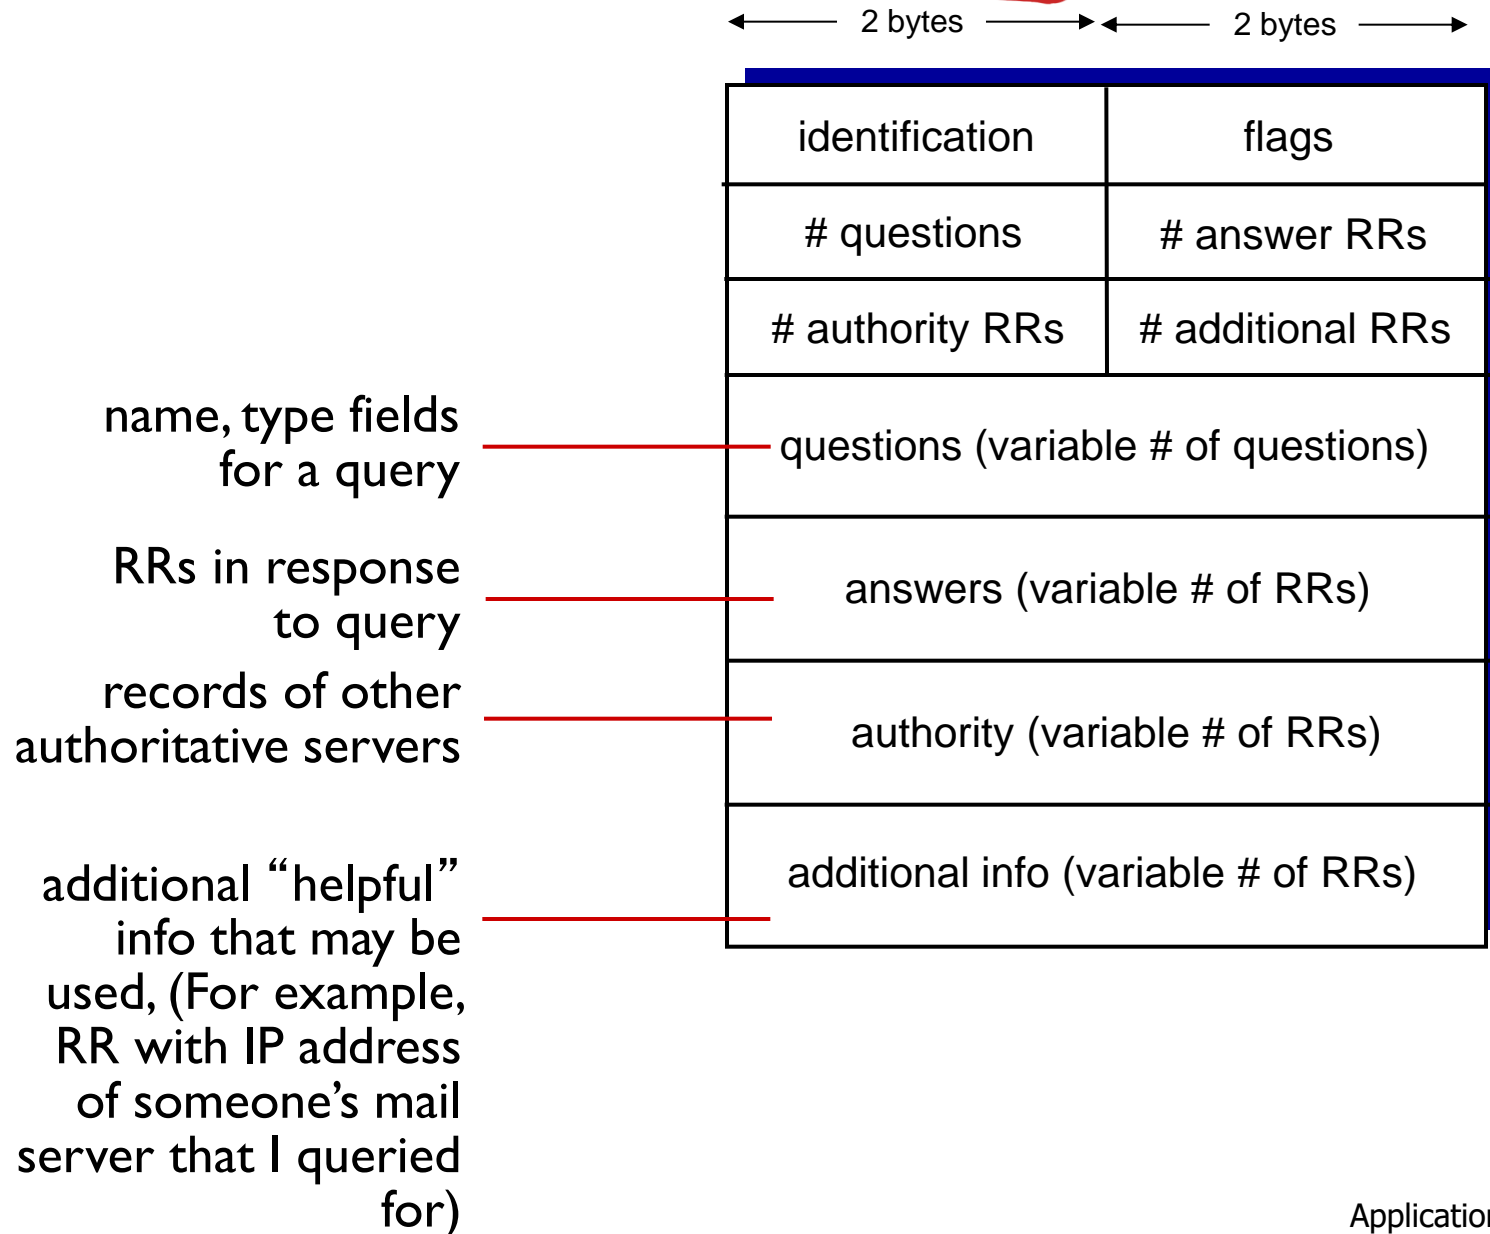
- *query* and *reply* messages, both with same *message format*

message header

- **identification**: 16 bit # for query, reply to query uses same #
- **flags**:
 - query(0) or reply(1)
 - recursion desired
 - recursion available
 - reply is authoritative (whether DNS Server is authoritative server. Sometimes there is intermediate DNS server which isn't authoritative)




DNS protocol, messages



How to insert Records into DNS?

- example: new startup “*Network Utopia*”
- **First step:** Assign names and IP addresses to your **web** server and **authoritative dns server**.



Inserting records into DNS

- example: new startup “*Network Utopia*”
- Register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
 - **2nd step:** provide names, IP addresses of authoritative name server (primary and secondary)
 - **3rd step:** registrar inserts two RRs into *.com TLD server*:
RR1: (*networkutopia.com*, *dns1.networkutopia.com*, **NS**)
RR 2: (*dns1.networkutopia.com*, *212.212.212.1*, **A**)
- **4th step:** In authoritative server add (Located on the Server premises)
 - type A record for *www.networkutopia.com*, e.g.,
 - *networkutopia.com*, *212.212.212.4*, **A**
 -

Attacking DNS

DDoS attacks

- bombard root servers with traffic
 - not successful to date
 - local DNS servers cache IPs of TLD servers, allowing root server bypass

redirect attacks

- man-in-middle
 - Intercept queries
- DNS poisoning
 - Send bogus replies to DNS server, which caches