

# Chapter 2

## Application Layer

### A note on the use of these Powerpoint slides:

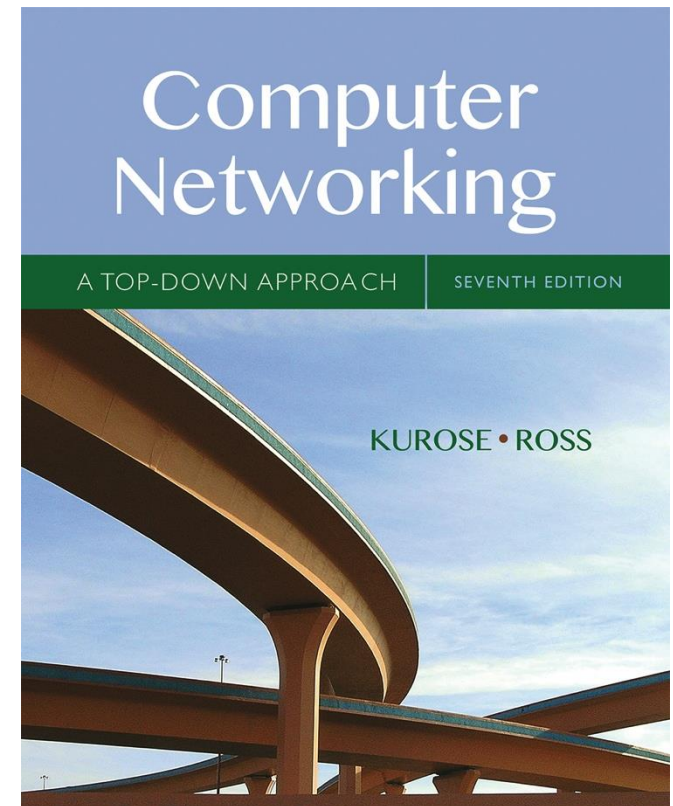
We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016

J.F Kurose and K.W. Ross, All Rights Reserved



## *Computer Networking: A Top Down Approach*

7<sup>th</sup> edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications


2.6 video streaming and content distribution networks

# Pure P2P architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

## *examples:*

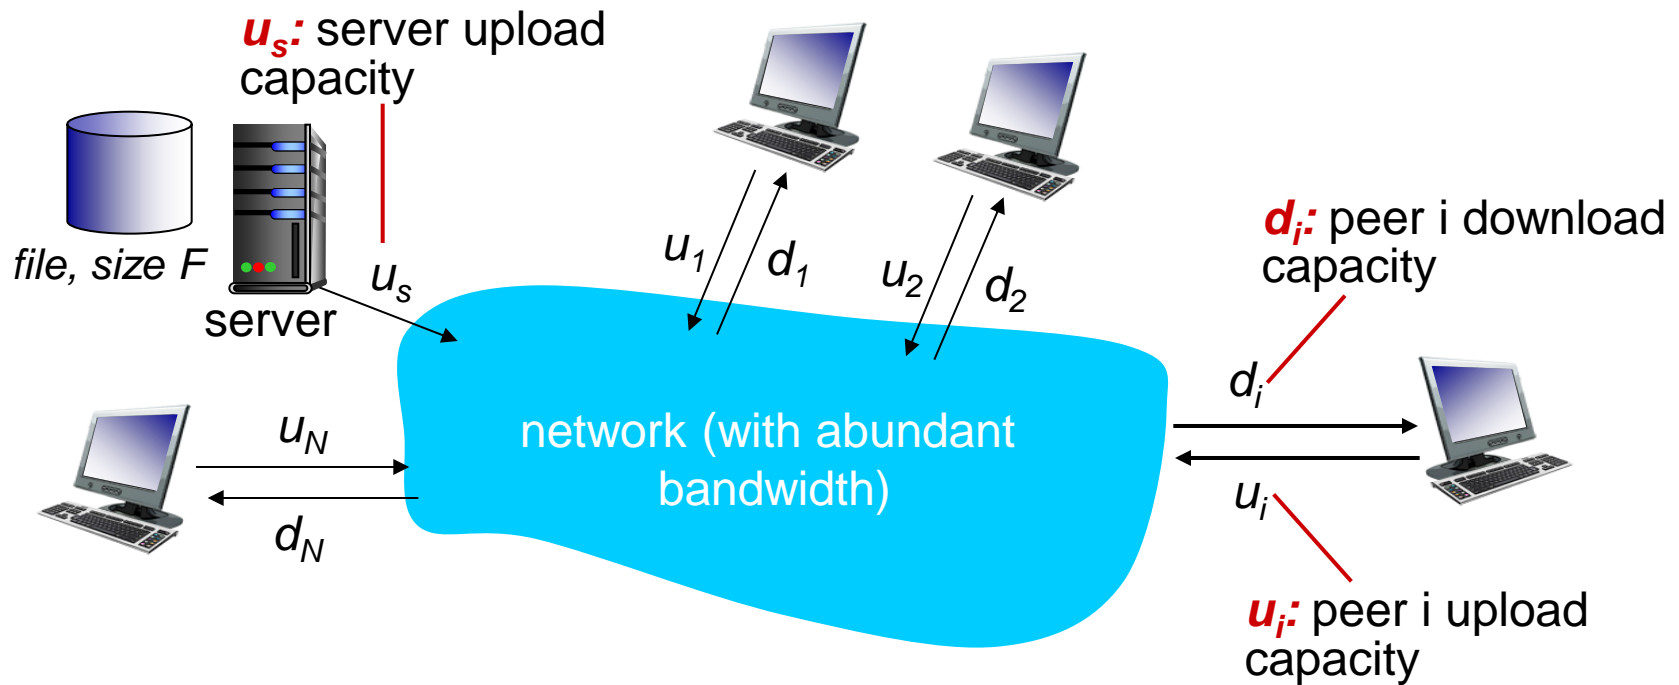
- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)



# File distribution: client-server vs P2P

Question: how much time to distribute file (size  $F$ ) from one server to  $N$  peers?


- peer upload/download capacity is limited resource



# File distribution time: client-server

- **server transmission:** must sequentially send (upload)  $N$  file copies:

- time to send one copy:  $F/u_s$
- time to send  $N$  copies:  $NF/u_s$



- **client:** each client must download file copy

- $d_{min}$  = min client download rate
- min client download time:  $F/d_{min}$


*time to distribute  $F$   
to  $N$  clients using  
client-server approach*

$$D_{c-s} \geq \max\{NF/u_s, F/d_{min}\}$$

increases linearly in  $N$

# File distribution time: P2P

- **server transmission:** must upload at least one copy
  - time to send one copy:  $F/u_s$
- **client:** each client must download file copy
  - min client download time:  $F/d_{\min}$
- **clients:** as aggregate must download  $NF$  bits
  - fastest possible upload rate (assuming all nodes sending file chunks to same peer):  $u_s + \sum_{i=1, N} u_i$



time to distribute  $F$  to  $N$   
clients using P2P  
approach

$$D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$

increases linearly in  $N$  ...

... but so does this, as each peer brings service capacity

# Client-server vs. P2P: example


All peers have the same upload rate, i.e.,  $u_1 = u_2 = \dots u_N = u$   
client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$

*For client-server*

$$D_{c-s} = NF/10u$$

*For P2P*

$$D_{P2P} = \max\{F/10u, NF/(N+10)u\}$$




# P2P file distribution: BitTorrent

- file divided into 256Kbytes chunks
- peers in torrent send/receive file chunks

**tracker:** tracks peers participating in torrent


**torrent:** group of peers exchanging chunks of a file





# P2P file distribution: BitTorrent

- peer joining torrent:
  - registers with tracker to get list of peers, connects to subset of peers (“neighbors”)
  - has no chunks, but will accumulate them over time from other peers
- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn**: peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent



# BitTorrent: requesting, sending file chunks

## *requesting chunks: rarest first*


- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests **missing chunks** from peers, the technique is called **rarest first**
  - Chunks having fewest repeating copies.

## *sending chunks: tit-for-tat*

- Alice sends chunks to those four peers currently sending her chunks **at highest rate**
  - other peers are choked by Alice (do not receive chunks from her)
  - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
  - “**optimistically unchoked**” this peer
  - newly chosen peer may join top 4

# BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks (CDNs)

# Internet video:

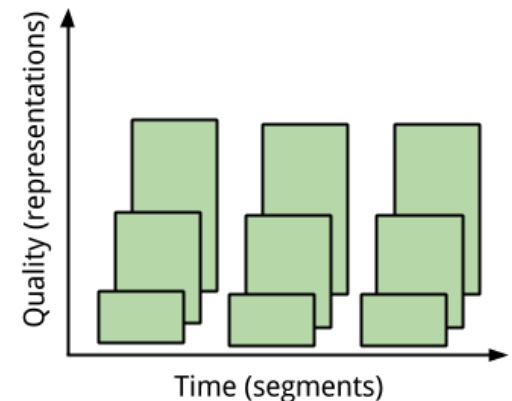
- Internet companies providing streaming videos are Netflix, Youtube(Google), Amazon, Youku
  - Netflix and Youtube, consumed 37% and 16% residential ISP traffic in 2015.
- A video is a sequence of images displayed at constant rate at 24 or 30 images per second.
- Videos can be compressed using compression algorithms.
- Trading off video quality with bit rate:
  - higher the bit rate → the better the image quality → better overall user viewing experience → needs more space for storage
  - Ranges from 100kbps for low quality to 3Mbps for streaming HD movies

# Streaming Multimedia using HTTP

- Using traditional HTTP streaming, a client requests a video from the server
  - Server sends video file in HTTP response message
  - Client collects bytes in application buffer and starts displaying them on screen
  - All clients receive the same encoding of the video regardless of the amount of bandwidth available to the clients.


# Streaming multimedia: DASH

- **DASH**: *D*ynamic, *A*daptive *S*treaming over *H*TTP
- *server*:
  - divides video file into multiple chunks
  - each chunk stored, encoded at different rates
  - *manifest file*: provides URLs for different chunks
- *client*:
  - periodically measures server-to-client bandwidth
  - consulting manifest, requests one chunk at a time
    - chooses maximum coding rate sustainable given current bandwidth
    - can choose different coding rates (and/or resolution) at different points in time (depending on available bandwidth at time)



# Streaming multimedia: DASH

- *DASH: Dynamic, Adaptive Streaming over HTTP*
- “intelligence” at client: client determines
  - *when* to request chunk (so that *buffer starvation*, or *overflow* does not occur)
  - *what encoding rate* to request (higher quality when more bandwidth available)
  - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)





# Video Streaming and CDNs: context

- video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- **Straight forward approach:** build a single massive data center, store all videos, stream videos to the clients directly
- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **option 1:** single, large “mega-server”
  - single point of failure
  - point of network congestion
  - long path to distant clients
  - multiple copies of video sent over outgoing link

YouTube

NETFLIX

hulu

 迅雷看看  
www.kankan.com  
—— 网络高清影院 ——

 Akamai

....quite simple but this solution *doesn't scale*

# Content distribution networks

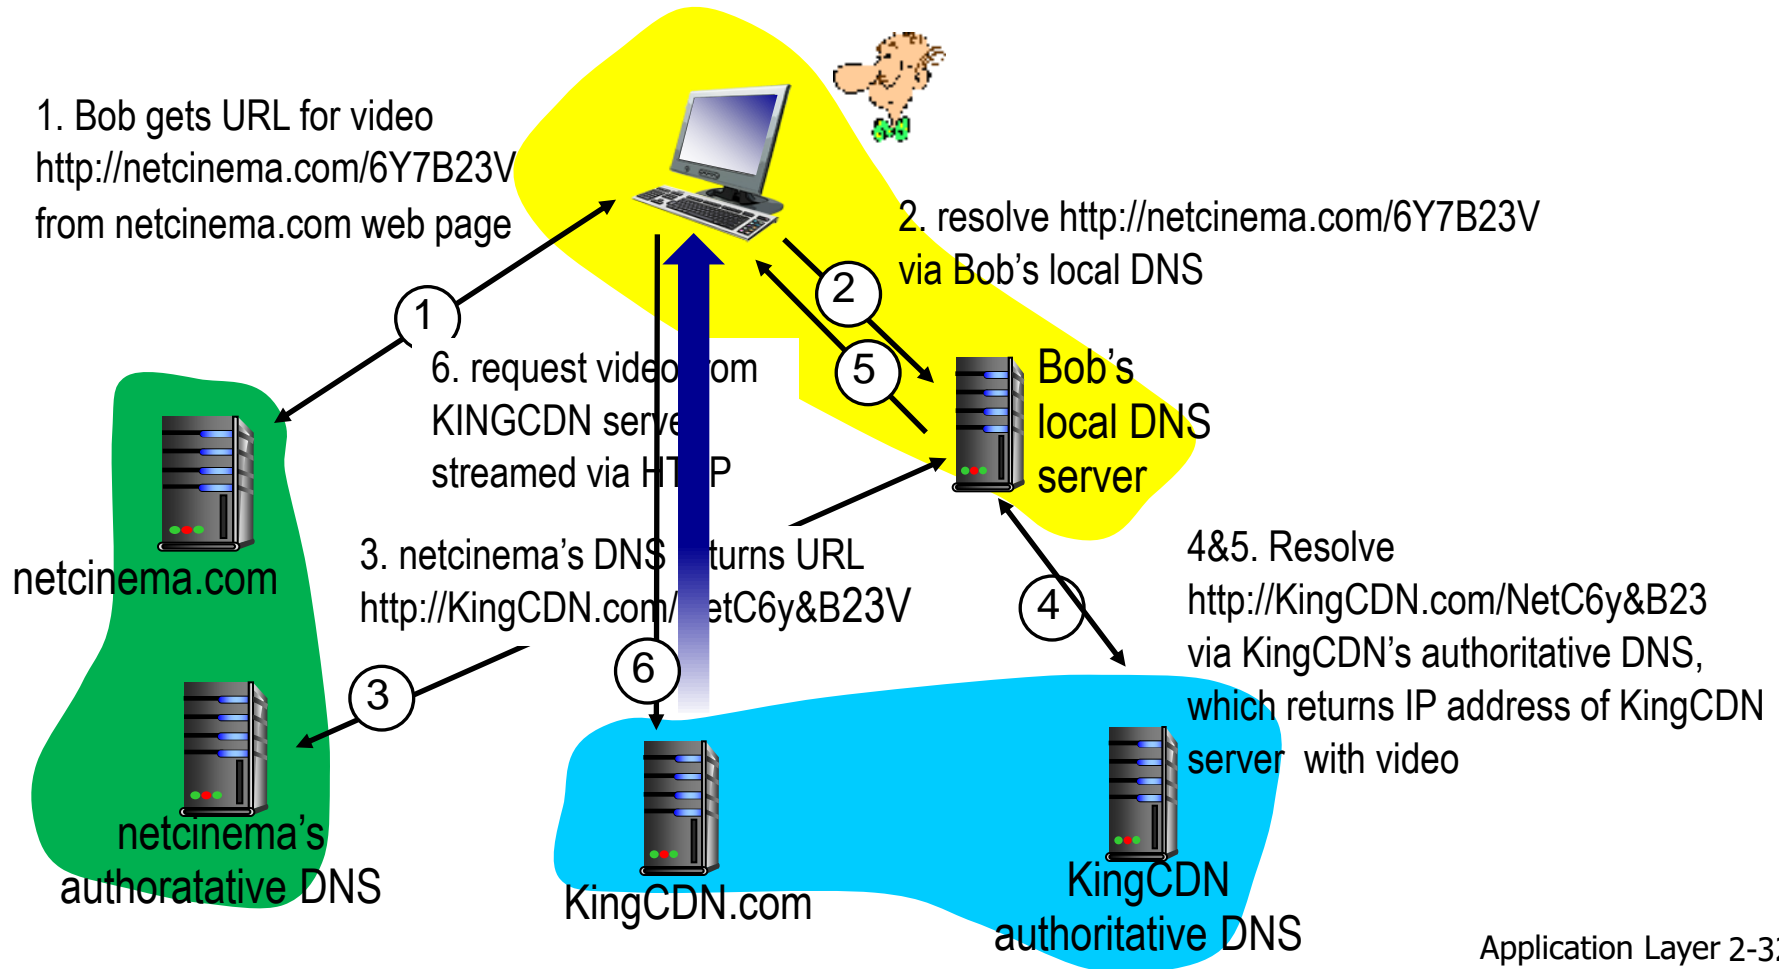
---

- *option 2*: store/serve multiple copies of videos at multiple geographically distributed sites (*CDN*)
  - *enter deep*: push CDN servers deep into many access networks
    - close to users, but the task of maintaining and managing the clusters becomes challenging
    - used by Akamai, 1700 locations
  - *bring home*: smaller number (10's) of larger clusters in POPs (Point of Presence) near (but not within) access networks
    - Lower maintenance and management overhead, but at the expense of higher delay and lower throughput to end users
    - used by Limelight

# CDN content access: a closer look

Bob (client) requests video **from provider Netcinema** <http://netcinema.com/6Y7B23V>

- video stored in 3<sup>rd</sup> party CDN company (**kingCDN**) at <http://KingCDN.com/NetC6y&B23V>
- CDN uses **intercepting and redirecting** the request.



# Cluster Selection Strategy

---

- ❖ *challenge*: how does CDN DNS select “good” CDN node to stream to client
  - pick CDN node geographically closest to client
  - pick CDN node with shortest delay (or min # hops) to client (CDN nodes periodically ping access ISPs, reporting results to CDN DNS)
- ❖ *alternative*: let *client* decide - give client a list of several CDN servers
  - client pings servers, picks “best”
  - Netflix approach

# **CASE STUDIES**


A thick, hand-drawn style red line underlining the text "CASE STUDIES".

# Case Study: Netflix

---

- ❖ owns very little infrastructure, uses 3<sup>rd</sup> party services:
  - own registration, payment servers
  - Amazon (3<sup>rd</sup> party) cloud services:
    - Netflix uploads studio master to Amazon cloud
    - create multiple version of movie (different encodings) in cloud
    - upload versions from cloud to CDNs
    - Cloud hosts Netflix web pages for user browsing
  - *three* 3<sup>rd</sup> party CDNs host/stream Netflix content: Akamai, Limelight, Level-3
- ❖ Netflix distributes by pushing the videos to its CDN servers during off-peak hours.
  - Pushes only the most popular videos, determined on day to day basis.

# Case study: Netflix



# Case study: Youtube

- ❖ *Private CDN*
- ❖ use DNS to redirect client to a cluster
  - usually smallest RTT between client and cluster
  - may be directed to more distant cluster for load balancing
  - may also be redirected if cluster doesn't have the file



# Case study: Kankan

- ❖ P2P video distribution, in China
- ❖ similar to BitTorrent
  - Client contacts tracker
  - download chunks of video from peers in parallel
  - focus on downloading chunks needed soon

# Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

# Chapter 2: summary

*most importantly: learned about protocols!*

- typical request/reply message exchange:
  - client requests info or service
  - server responds with data, status code
- message formats:
  - *headers*: fields giving info about data
  - *data*: info(payload) being communicated

*important themes:*

- control vs. messages
- centralized vs. decentralized
- reliable vs. unreliable message transfer
- “complexity at network edge”