

# Chapter 3

## Transport Layer

### Part 4/5

---

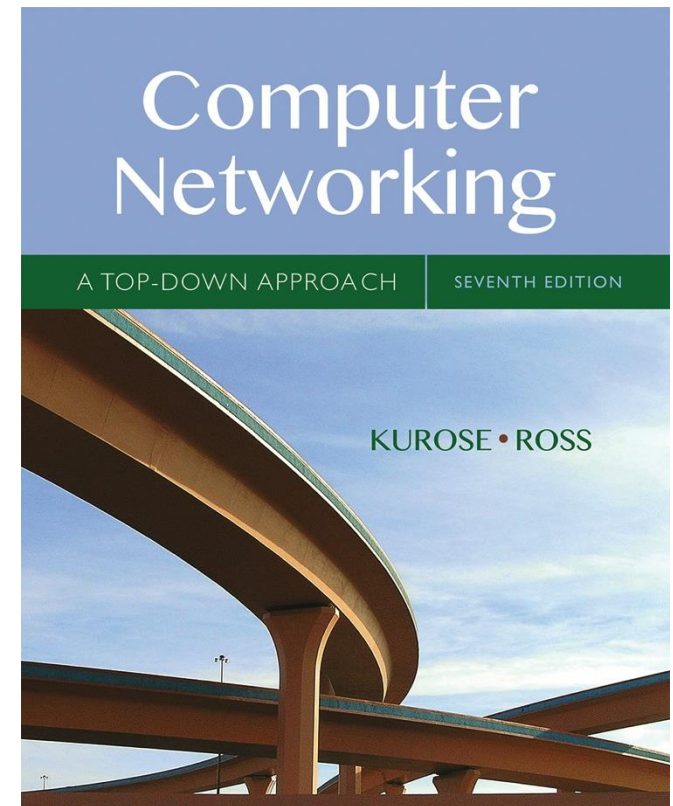
#### A note on the use of these Powerpoint slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

© All material copyright 1996-2016  
J.F Kurose and K.W. Ross, All Rights Reserved



## *Computer Networking: A Top Down Approach*

7<sup>th</sup> edition

Jim Kurose, Keith Ross

Pearson/Addison Wesley

April 2016

# Chapter 3 outline

3.1 transport-layer services

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

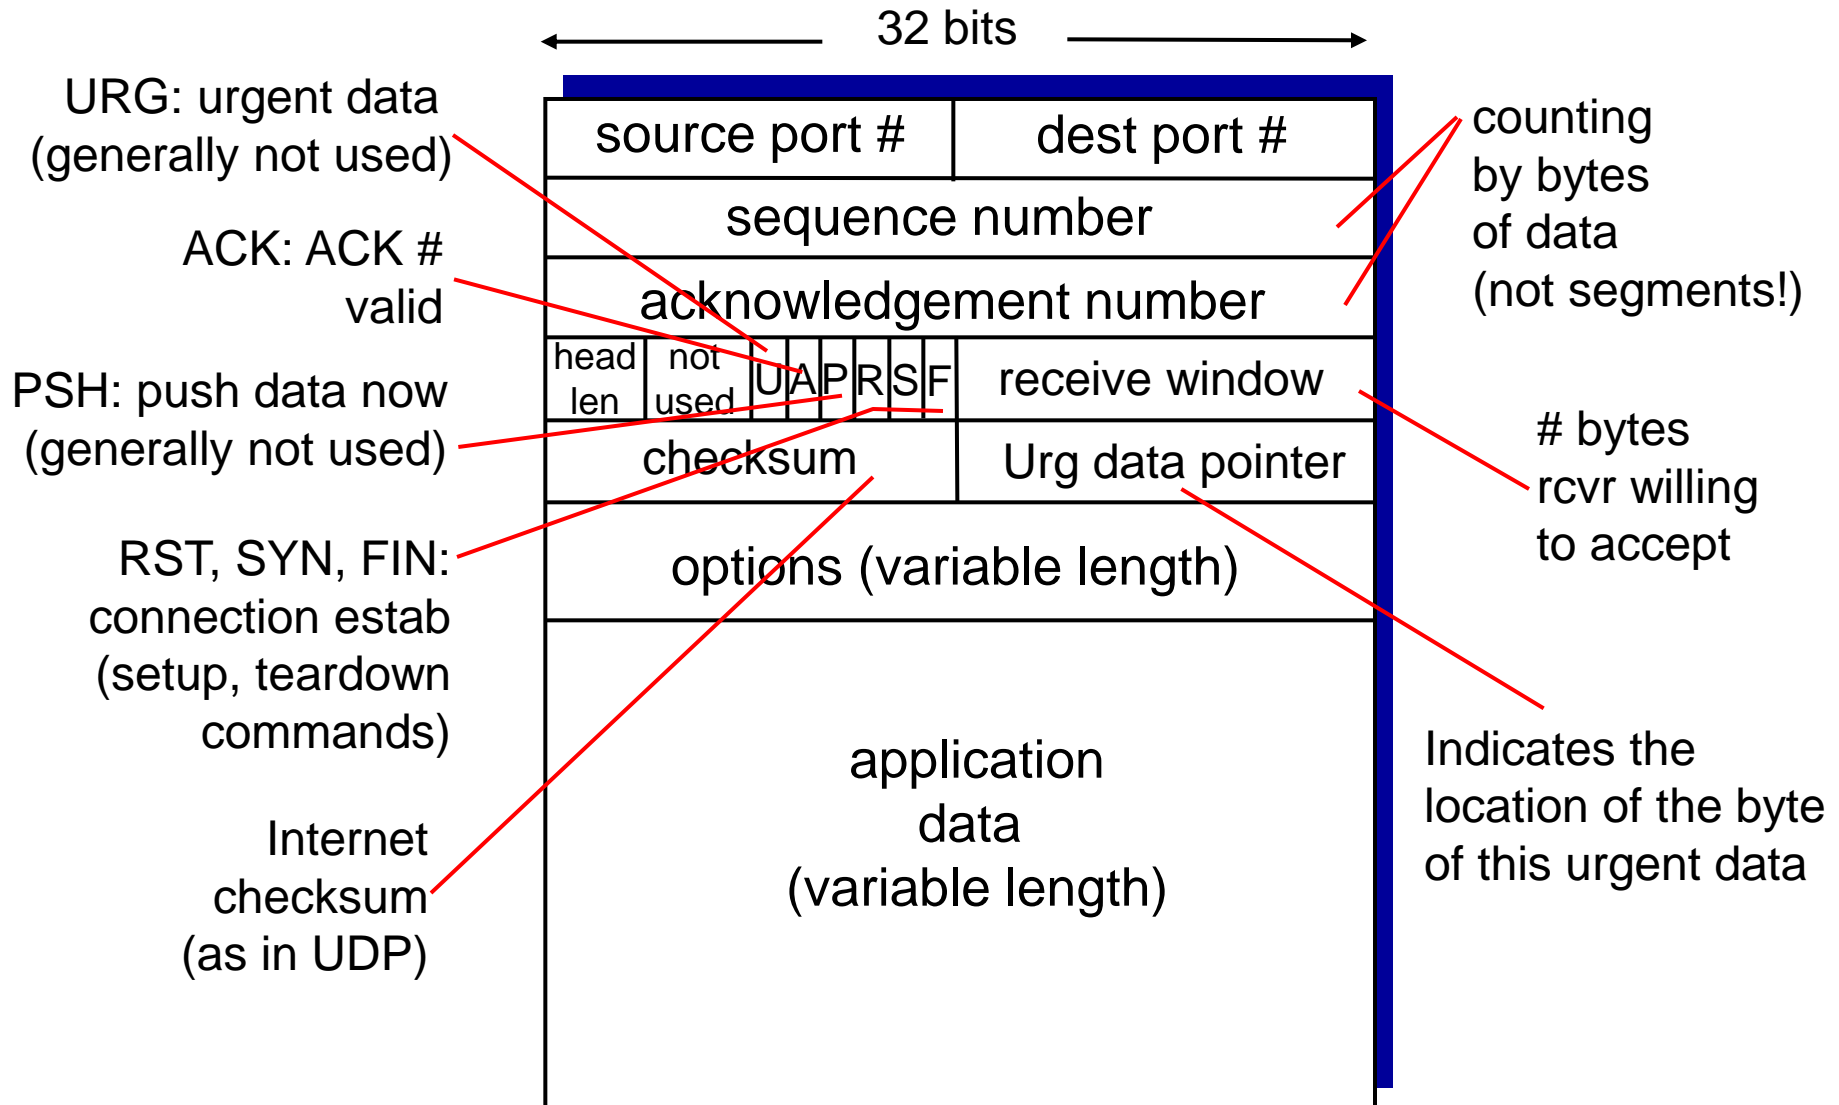
3.7 TCP congestion control

# TCP: Overview

RFCs: 793, 1122, 1323, 2018, 2581

- **point-to-point:**
  - one sender, one receiver
- **reliable, in-order *byte stream*:**
  - Segments delivered to app in order
- **pipelined:**
  - TCP congestion and flow control set window size
- **full duplex data:**
  - bi-directional data flow in same connection
  - MSS: maximum segment size
- **connection-oriented:**
  - handshaking (exchange of control msgs) initiates sender, receiver state before data exchange
- **flow controlled:**
  - sender will not overwhelm receiver

# TCP segment structure



# TCP seq. numbers, ACKs

## sequence numbers:

- byte stream “number” of first byte in segment’s data

## acknowledgements:

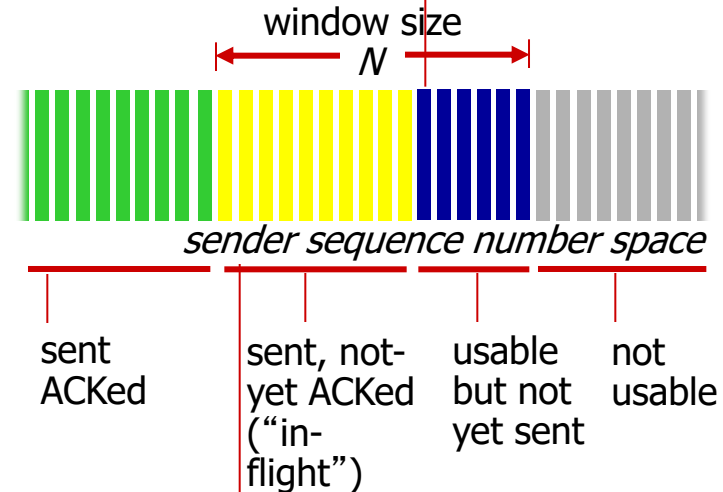
- seq # of next byte expected from other side
- cumulative ACK

**Q:** how receiver handles out-of-order segments

- **A:** TCP spec doesn’t say,  
- up to implementor

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	A
checksum	urg pointer

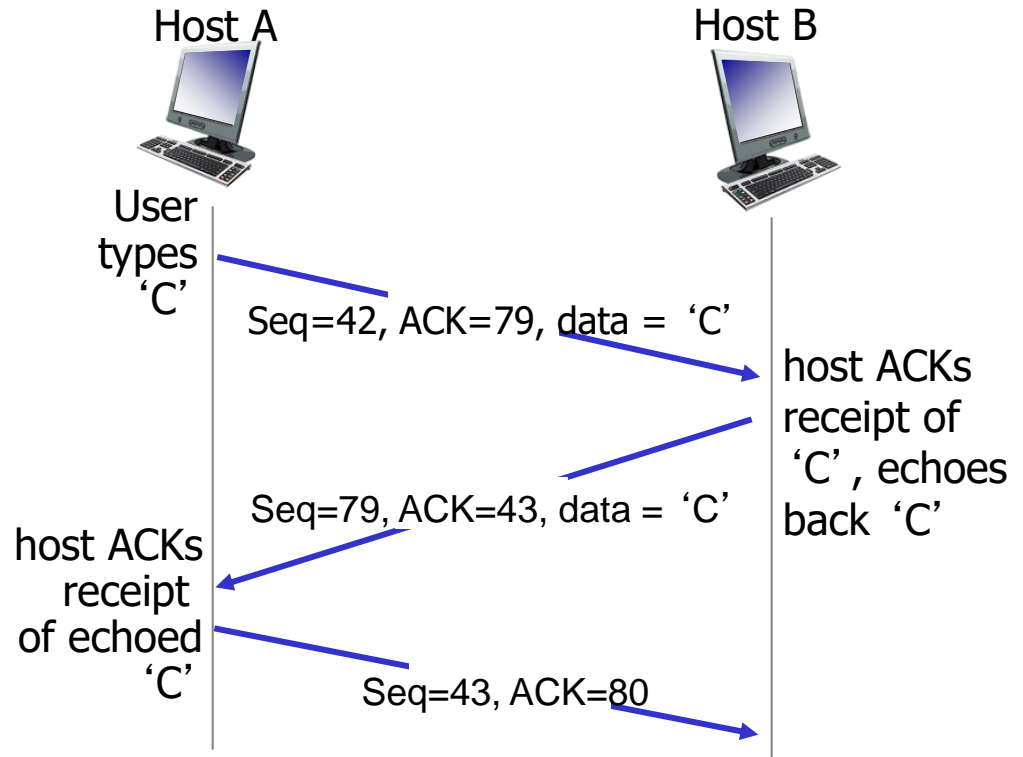
# TCP seq. numbers, ACKs

## Piggybacked ACK:

- ACK for client-to-server data is carried in a segment carrying server-to-client data.

## acknowledgements:

- ACK=43 in server-to-client segment shows server has received everything upto byte 42
- Server waiting for byte 43



# TCP round trip time and timeout

## Why need to estimate RTT?

- ❑ "timeout" and "retransmit" needed to address packet loss
- ❑ need to know when to timeout and retransmit

## Ideal world:

- ❑ exact RTT is needed

## Real world:

- ❑ RTTs change over time:
  - ❖ packets may take different paths
  - ❖ network load changes over time
- ❑ RTTs can only be estimated

## Some intuition

- ❑ What happens if too short?
  - ❖ Premature timeout
  - ❖ Unnecessary retransmissions
- ❑ What happens if too long?
  - ❖ Slow reaction to segment loss

# TCP round trip time and timeout

Q: how to estimate RTT?

- **SampleRTT**: measured time from segment transmission until ACK receipt
  - ignore retransmissions
- **SampleRTT** will vary, want estimated RTT “smoother”
  - average several *recent* measurements, not just current

**SampleRTT**




# TCP round trip time and timeout

- We also want to give more recent measurements **higher weight** in case things do change

$$\text{EstimatedRTT}^{(\text{current})} = (1 - \alpha) * \text{EstimatedRTT}^{(\text{Previous})} + \alpha * \text{SampleRTT}^{(\text{recent})}$$

- New EstimatedRTT is weighted combination of previous EstimatedRTT and new SampleRTT
- All the samples are averaged using **exponential weighted moving average** (EWMA)
- influence of past sample decreases exponentially fast
- Coefficient  $\alpha$  is the **degree of weighting decrease**
- $0 < \alpha < 1$ , but typical value:  $\alpha = 0.125$

# TCP round trip time and timeout



# TCP round trip time and timeout

## Setting the timeout

- ❑ `timeout = EstimatedRTT`, any problem with this???
- ❑ add a "safety margin" to `EstimatedRTT`
  - ❖ large variation in `EstimatedRTT` -> larger safety margin
- ❑ see how much `SampleRTT` deviates from `EstimatedRTT`:

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically,  $\beta = 0.25$ )

Then set timeout interval:

$$\text{TimeoutInterval} = \text{EstimatedRTT} + 4 * \text{DevRTT}$$

# TCP round trip time and timeout

- **timeout interval:** **EstimatedRTT** plus “safety margin”
  - large variation in **EstimatedRTT** → larger safety margin
- estimate **SampleRTT** deviation from **EstimatedRTT**:

$$\text{SampleDevRTT}^{(\text{recent})} = |\text{SampleRTT} - \text{EstimatedRTT}|$$

$$\text{DevRTT}^{(\text{current})} = (1 - \beta) * \text{DevRTT}^{(\text{Previous})} + \beta * \text{SampleDevRTT}^{(\text{recent})}$$

$(0 < \beta < 1; \text{ typically } \beta = 0.25)$

$$\text{TimeoutInterval} = \text{EstimatedRTT}^{(\text{current})} + 4 * \text{DevRTT}^{(\text{Current})}$$



↑  
estimated RTT

↑  
“safety margin”

# TCP round trip time and timeout (class exercises)

- Suppose that TCP's current estimatedRTT and DevRTT are 340 msec and 18 msec, respectively. Suppose that the next three SampleRTTs are 400, 270, and 390 respectively. Compute TCP's new value of estimatedRTT, DevRTT, and the TCP timeout value after each of these three measured RTT values is obtained. Use the values of  $\alpha = 0.125$  and  $\beta = 0.25$ .

$$\text{EstimatedRTT}_{(\text{current})} = (1 - \alpha) * \text{EstimatedRTT}_{(\text{Previous})} + \alpha * \text{SampleRTT}_{(\text{recent})}$$

$$\text{DevRTT}_{(\text{current})} = (1 - \beta) * \text{DevRTT}_{(\text{Previous})} + \beta * (|\text{SampleRTT} - \text{EstimatedRTT}|)_{(\text{recent})}$$

- Solution:

1)  $\text{SampleRTT}_{1(\text{recent})} = 400\text{ms}$ ,  $\text{EstimatedRTT}_{(\text{previous})} = 340\text{ms}$ ,  $\text{DevRTT}_{(\text{previous})} = 18\text{ms}$

$$\text{estimatedRTT}_{(\text{current})} = ?$$

$$\text{DevRTT}_{(\text{current})} = ?$$

$$\text{TimeoutInterval} = ?$$

# TCP round trip time and timeout (class exercises)

- Suppose that TCP's current estimatedRTT and DevRTT are 340 msec and 18 msec, respectively. Suppose that the next three SampleRTTs are 400, 270, and 390 respectively. Compute TCP's new value of estimatedRTT, DevRTT, and the TCP timeout value after each of these three measured RTT values is obtained. Use the values of  $\alpha = 0.125$  and  $\beta = 0.25$ .

$$\text{EstimatedRTT}^{(\text{current})} = (1 - \alpha) * \text{EstimatedRTT}^{(\text{Previous})} + \alpha * \text{SampleRTT}^{(\text{recent})}$$

$$\text{DevRTT}^{(\text{current})} = (1 - \beta) * \text{DevRTT}^{(\text{Previous})} + \beta * (|\text{SampleRTT} - \text{EstimatedRTT}|)^{(\text{recent})}$$

- Solution:

**1) SampleRTT<sub>1</sub> = 400ms, EstimatedRTT=340ms, DevRTT=18ms**

$$\text{estimatedRTT} = 0.875 * 340 + 0.125 * 400 = 347.5 \text{ msec}$$

$$\text{DevRTT} = 0.75 * 18 + 0.25 * (\text{abs}(400 - 340)) = 28.5 \text{ msec}$$

$$\text{TimeoutInterval} = 347.5 + 4 * 28.5 = 461.5 \text{ msec}$$

# TCP round trip time and timeout (class exercises)

**2) SampleRTT<sub>2</sub> = 270ms , EstimatedRTT=347.5ms, DevRTT=28.5ms**

$$\text{estimatedRTT} = 0.875 * 347.5 + 0.125 * 270 = 337.8125 \text{ msec}$$

$$\text{DevRTT} = 0.75 * 28.5 + 0.25 * (\text{abs}(270 - 347.5)) = 40.75 \text{ msec}$$

$$\text{TimeoutInterval} = 337.8125 + 4 * 40.625 = ? \text{ msec}$$

**3) SampleRTT<sub>3</sub> = 390ms , EstimatedRTT=337.8125ms, DevRTT=40.625ms**

$$\text{estimatedRTT} = 0.875 * 337.8125 + 0.125 * 390 = 344.335 \text{ msec}$$

$$\text{DevRTT} = ??$$

$$\text{TimeoutInterval} = ??$$

# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control



# TCP reliable data transfer

- TCP creates rdt service on top of IP's unreliable service
  - pipelined segments
  - cumulative acks
- retransmissions triggered by:
  - timeout events
  - duplicate acks

let's initially consider simplified TCP sender:

- ignore duplicate acks
- ignore flow control
- ignore congestion control

# TCP sender events:

## *data rcvd from app:*

- create segment with seq #
- seq # is byte-stream number of first data byte in segment
- start timer if not already running
  - think of timer as for oldest unacked segment
  - expiration interval: `TimeoutInterval`

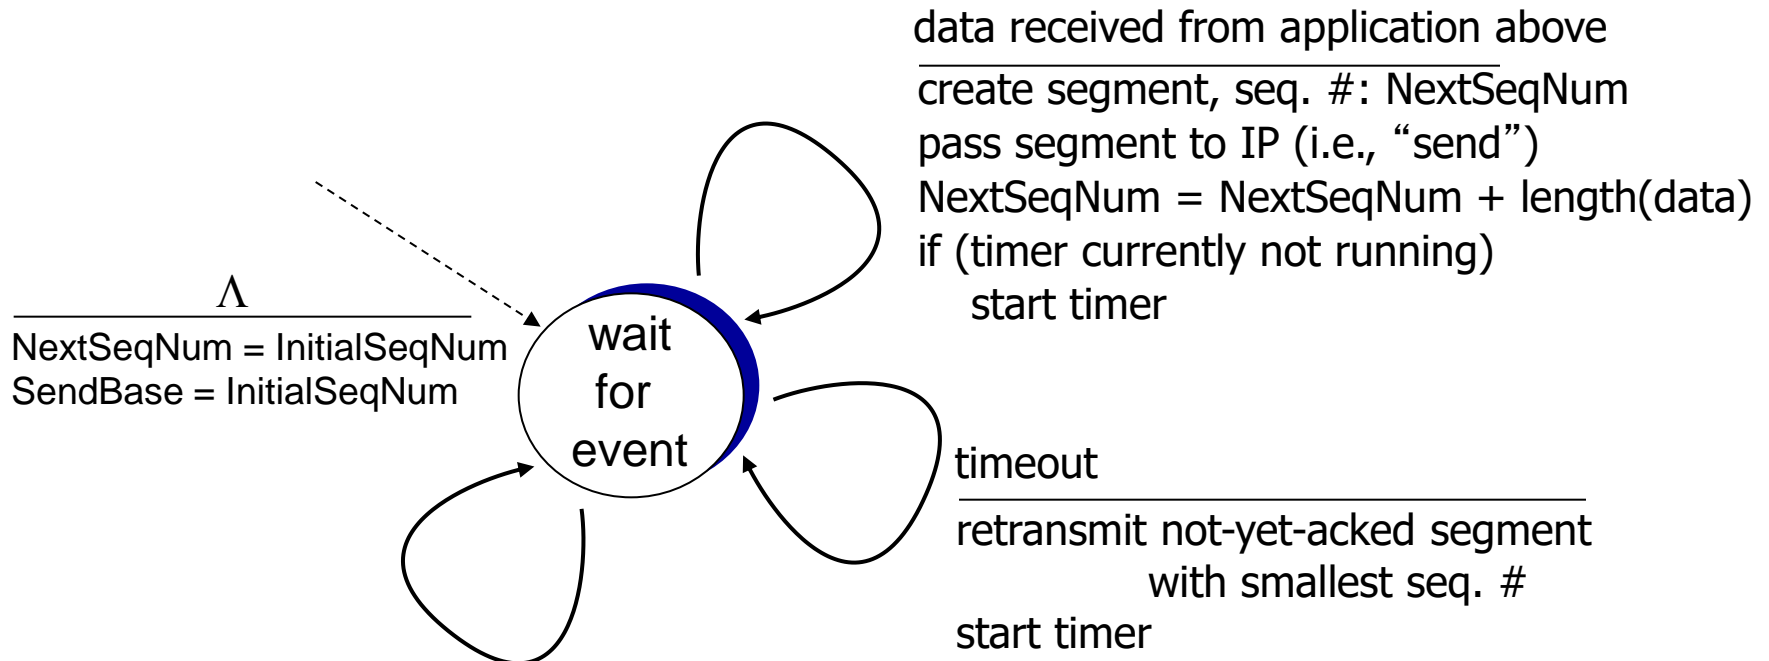
## *timeout:*

- retransmit segment that caused timeout
- restart timer

## *ack rcvd:*

- if ack acknowledges previously unacked segments
  - update what is known to be ACKed
  - start timer if there are still unacked segments


# TCP sender (simplified)



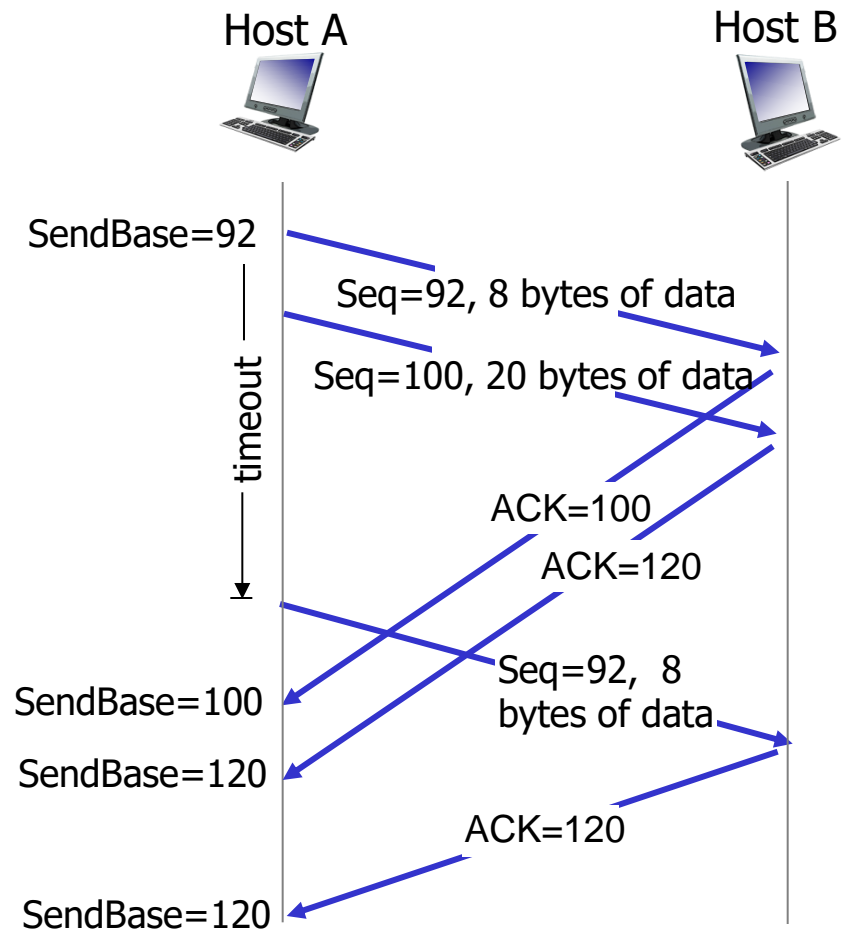
ACK received, with ACK field value y

```
if (y > SendBase) {  
    SendBase = y /* update SendBase */  
    /* SendBase-1: last cumulatively ACKed byte */  
    if (there are currently not-yet-acked segments)  
        start timer  
    else stop timer //no in-flight packet, so no need to start timer  
}
```

# TCP: retransmission scenarios

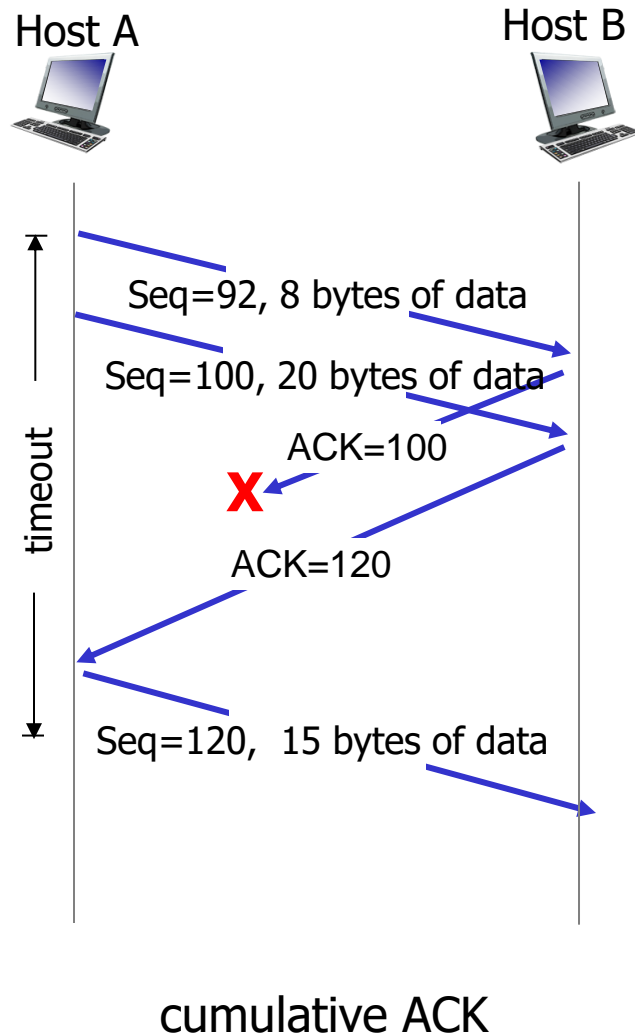


lost ACK scenario



premature timeout  
(In the last, receiver sends Cumulative ACK)  
Transport Layer 3-29

# TCP: retransmission scenarios



# TCP fast retransmit

- time-out period often relatively long:
  - long delay before resending lost packet
- detect lost segments via duplicate ACKs.
  - sender often sends many segments back-to-back
  - if segment is lost, there will likely be many duplicate ACKs.

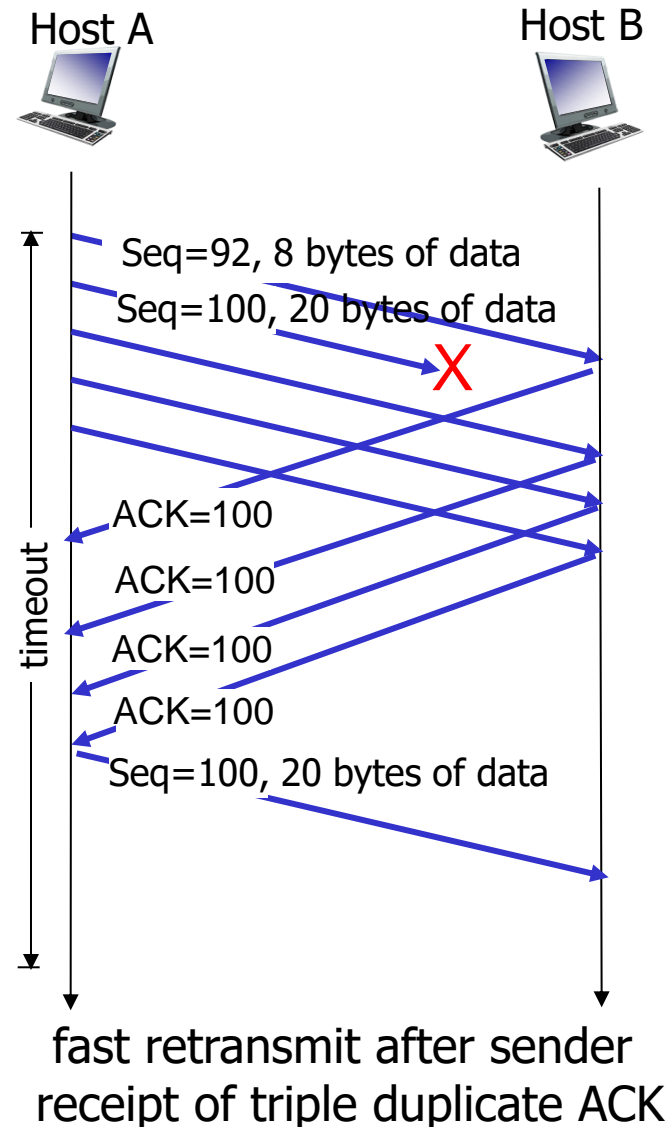
## *TCP fast retransmit*

if sender receives 3 ACKs for same data (“triple duplicate ACKs”), resend unacked segment with smallest seq #

- likely that unacked segment lost, so don't wait for timeout

# TCP fast retransmit

- Normally TCP sender will wait for the time-out before retransmission,
  - But in TCP fast retransmit, sender will retransmit before the timeout happens if it receives 3 dup ack.



# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control

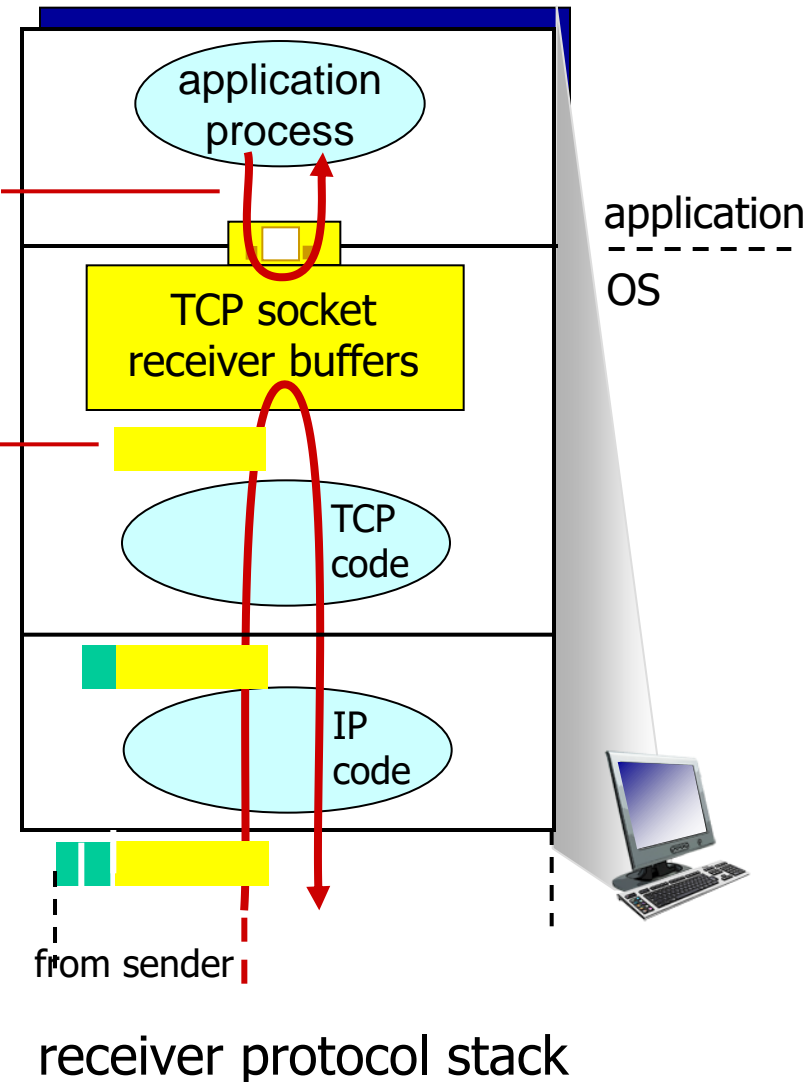


# TCP flow control

application may  
remove data from  
TCP socket buffers ....

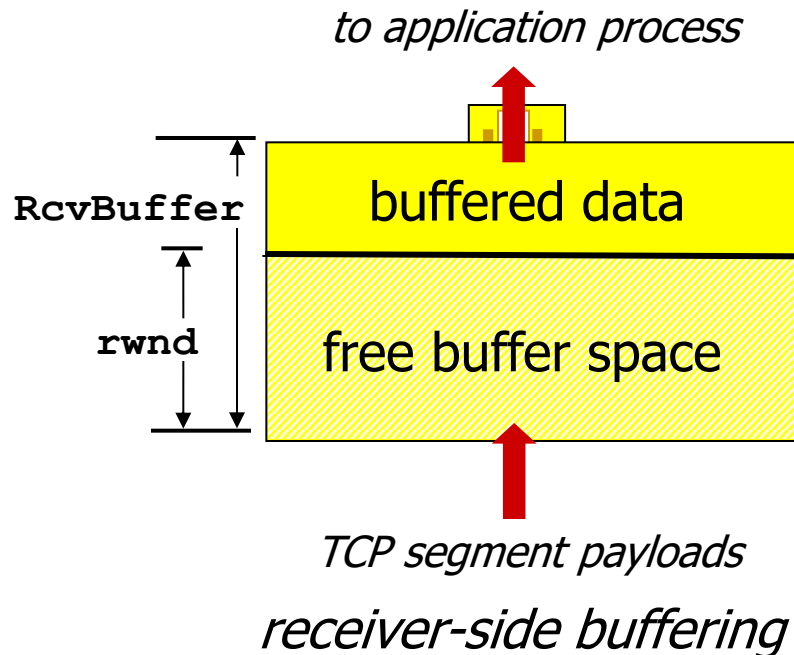
... slower than TCP  
receiver is delivering  
(sender is sending)

*flow control*  
receiver controls sender, so  
sender won't overflow  
receiver's buffer by transmitting  
too much, too fast



# TCP flow control

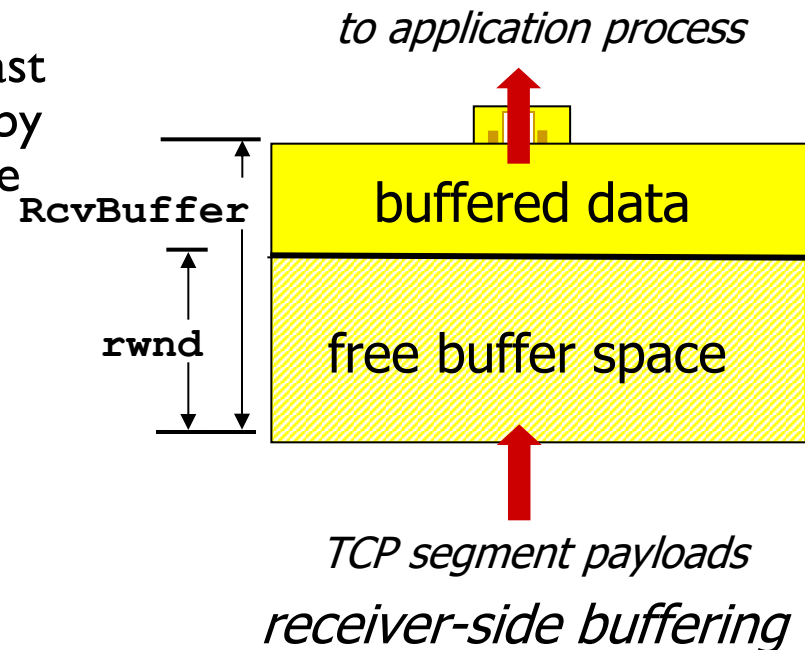
- receiver “advertises” free buffer space by including **rwnd** value in TCP header of receiver-to-sender segments
  - **RcvBuffer** size set via socket options (typical default is 4096 bytes)
  - many operating systems autoadjust **RcvBuffer**
- sender limits amount of unacked (“in-flight”) data to receiver’s **rwnd** value
- Flow control guarantees receive buffer will not overflow



# TCP flow control: how to calculate rwnd?

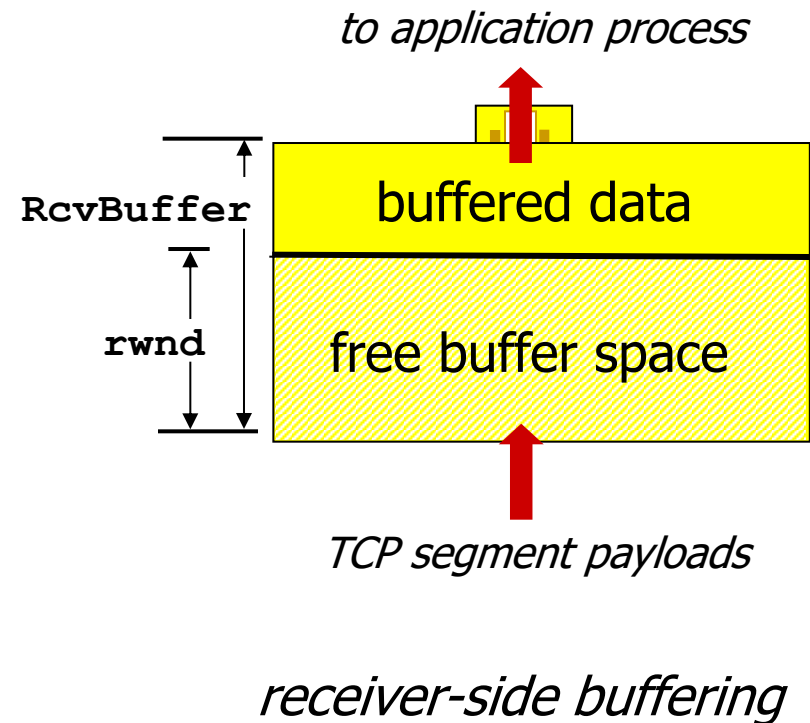
- Calculating rwnd
  - $\text{rwnd} = \text{RcvBuffer} - [\text{LastByteRcvd} - \text{LastByteRead}]$
- **LastByteRead:** the **sequence number** of the last byte in the data stream read from the buffer by the **application process** of the receiving host
- **LastByteRcvd:** the **sequence number** of the last byte in the data stream that has been received by the receive buffer of the receiving host from the **network**
- Relationship between RcvBuffer, LastByteRcvd, LastByteRead

$$\text{LastByteRcvd} - \text{LastByteRead} \leq \text{RcvBuffer}$$



# TCP flow control

- Deadlock Situation
  - Receiver consumed some data and free up some receive buffer space
  - Sender thinks  $rwnd=0$  and doesn't send segment
  - Sender has some data in the *send buffer*, it *periodically* sends a *one byte TCP segment* to the receiver to *trigger a response* from the receiver
  - **Ack** for the **byte size probe TCP segment** will contain the **non-zero value  $rwnd$**  that the **sender** uses for transmission



# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management


3.6 principles of congestion control

3.7 TCP congestion control


# Connection Management

before exchanging data, sender/receiver “handshake”:

- agree to establish connection (each knowing the other willing to establish connection)
- agree on connection parameters



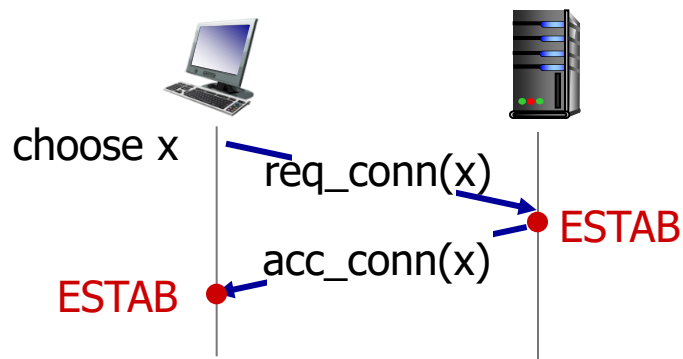

```
Socket clientSocket =  
    newSocket("hostname", "port  
    number");
```



```
Socket connectionSocket =  
    welcomeSocket.accept();
```

# Connection Management: Agreeing to establish a connection

2-way handshake:

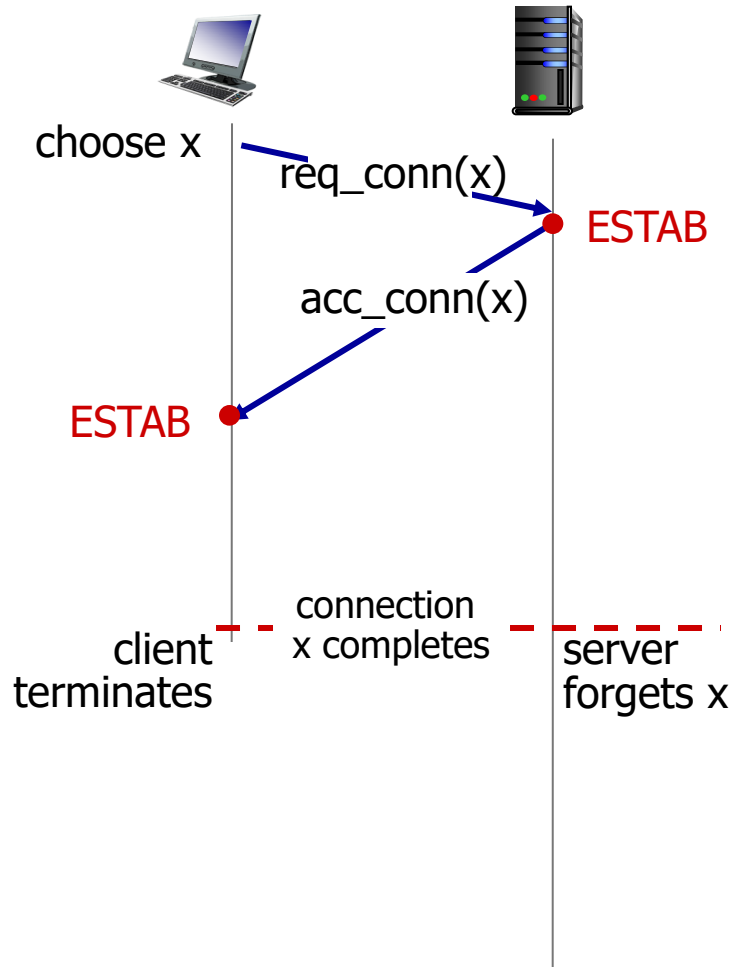


Q: will 2-way handshake always work in network?

- variable delays
- retransmitted messages (e.g. req\_conn(x)) due to message loss
- message reordering
- can't "see" other side

# Agreeing to establish a connection: 1st scenario

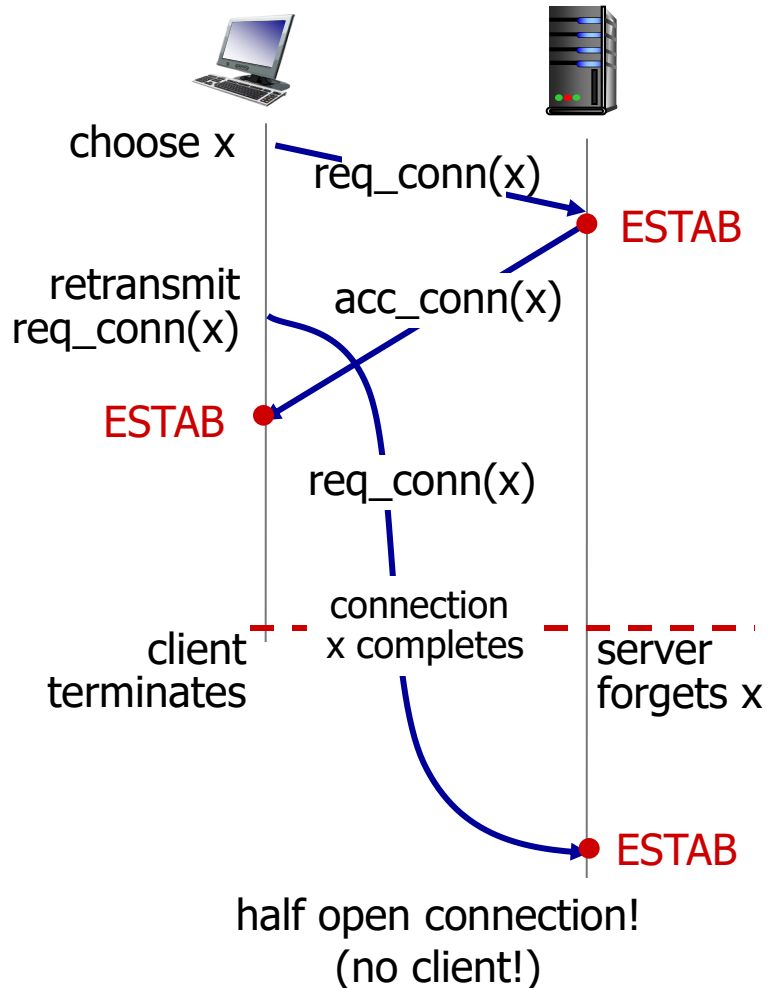
2-way handshake failure scenarios:





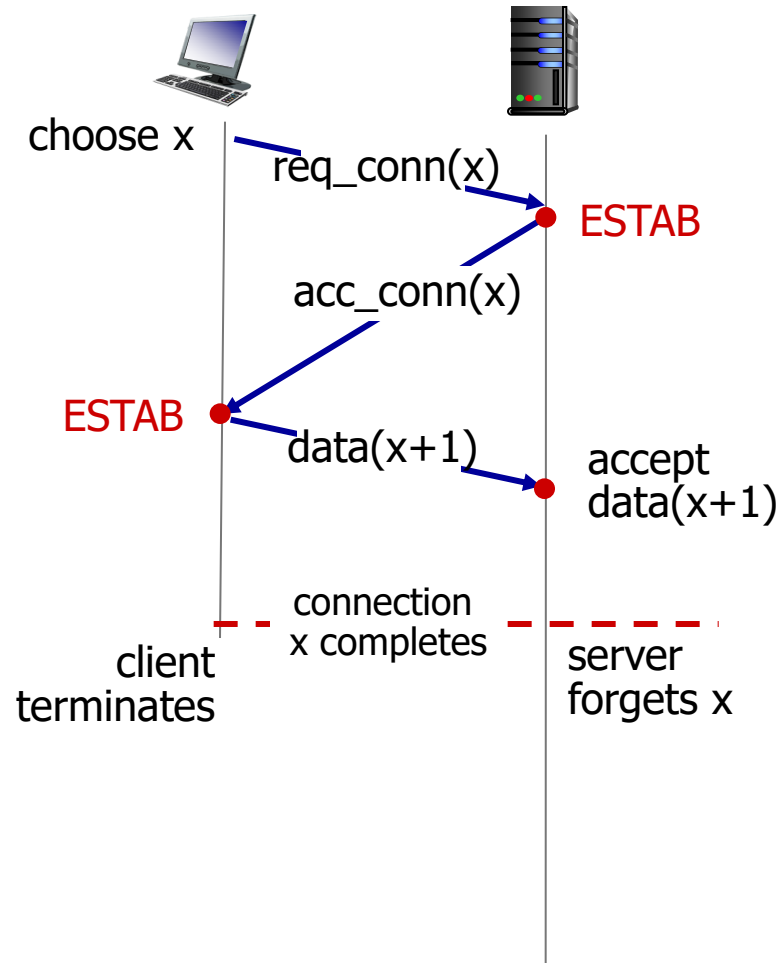
# Agreeing to establish a connection: 1st Scenario

2-way handshake failure scenarios:



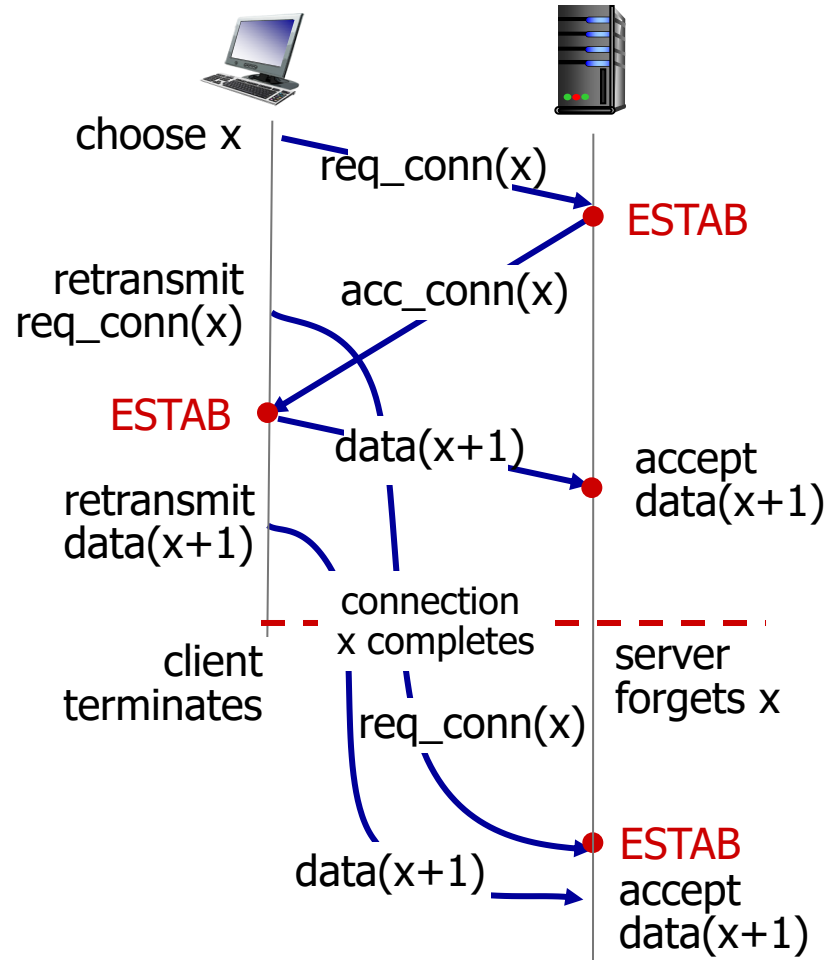
# Agreeing to establish a connection: 2<sup>nd</sup> Scenario

2-way handshake failure scenarios:

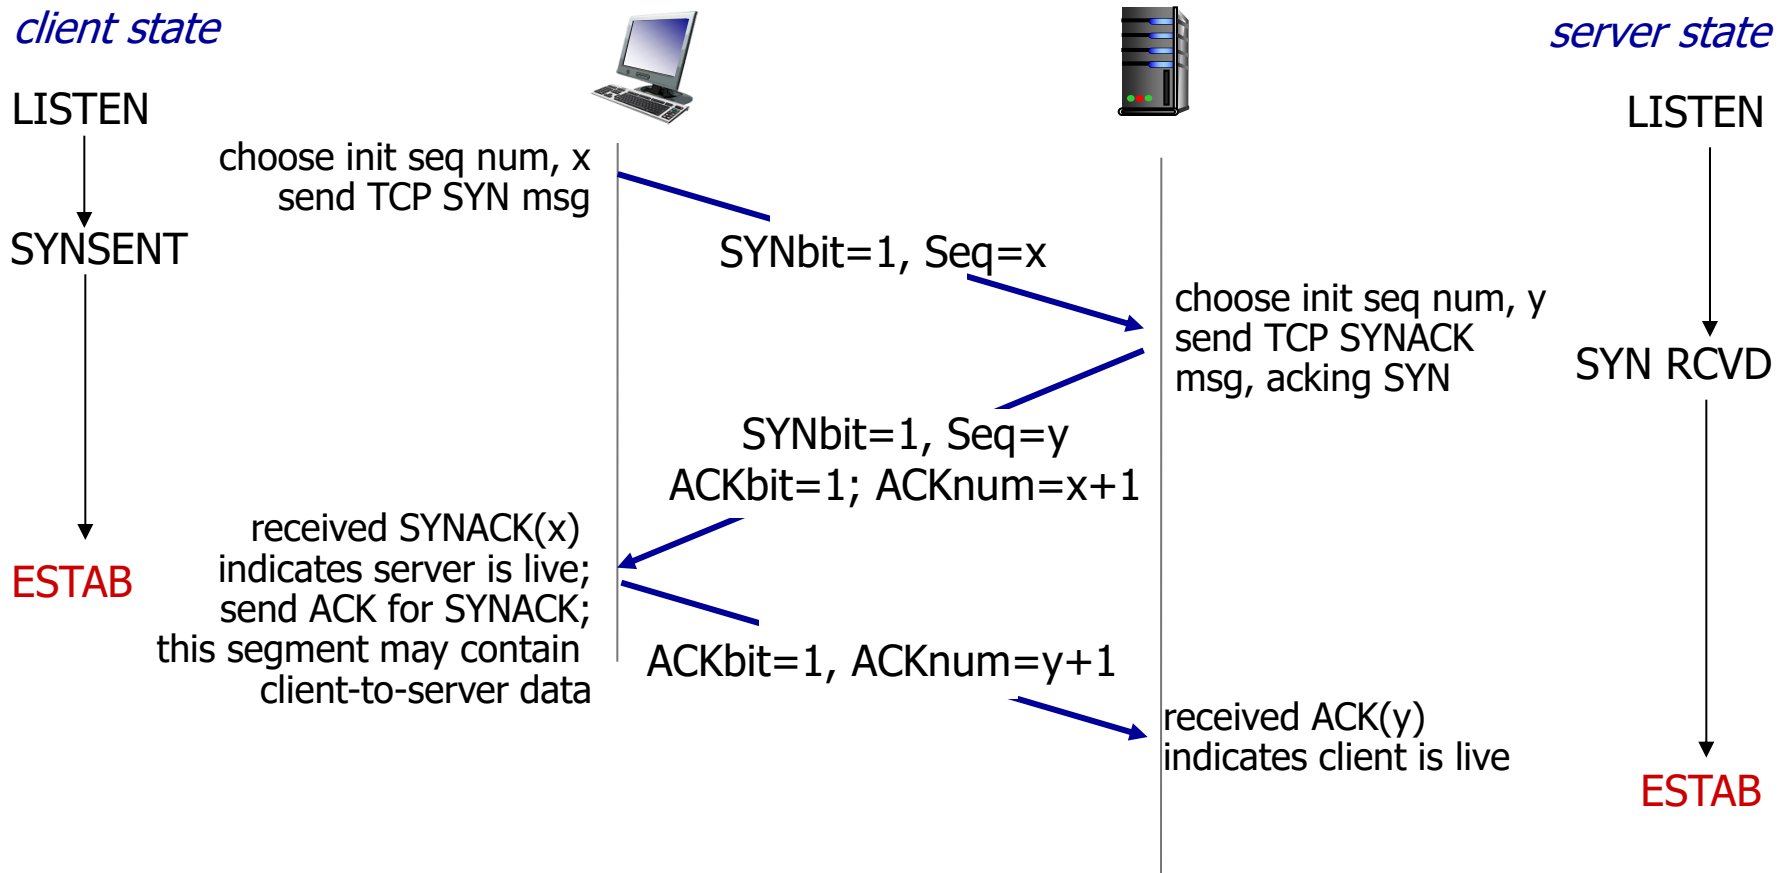


# Agreeing to establish a connection: 2<sup>nd</sup> Scenario

2-way handshake failure scenarios:

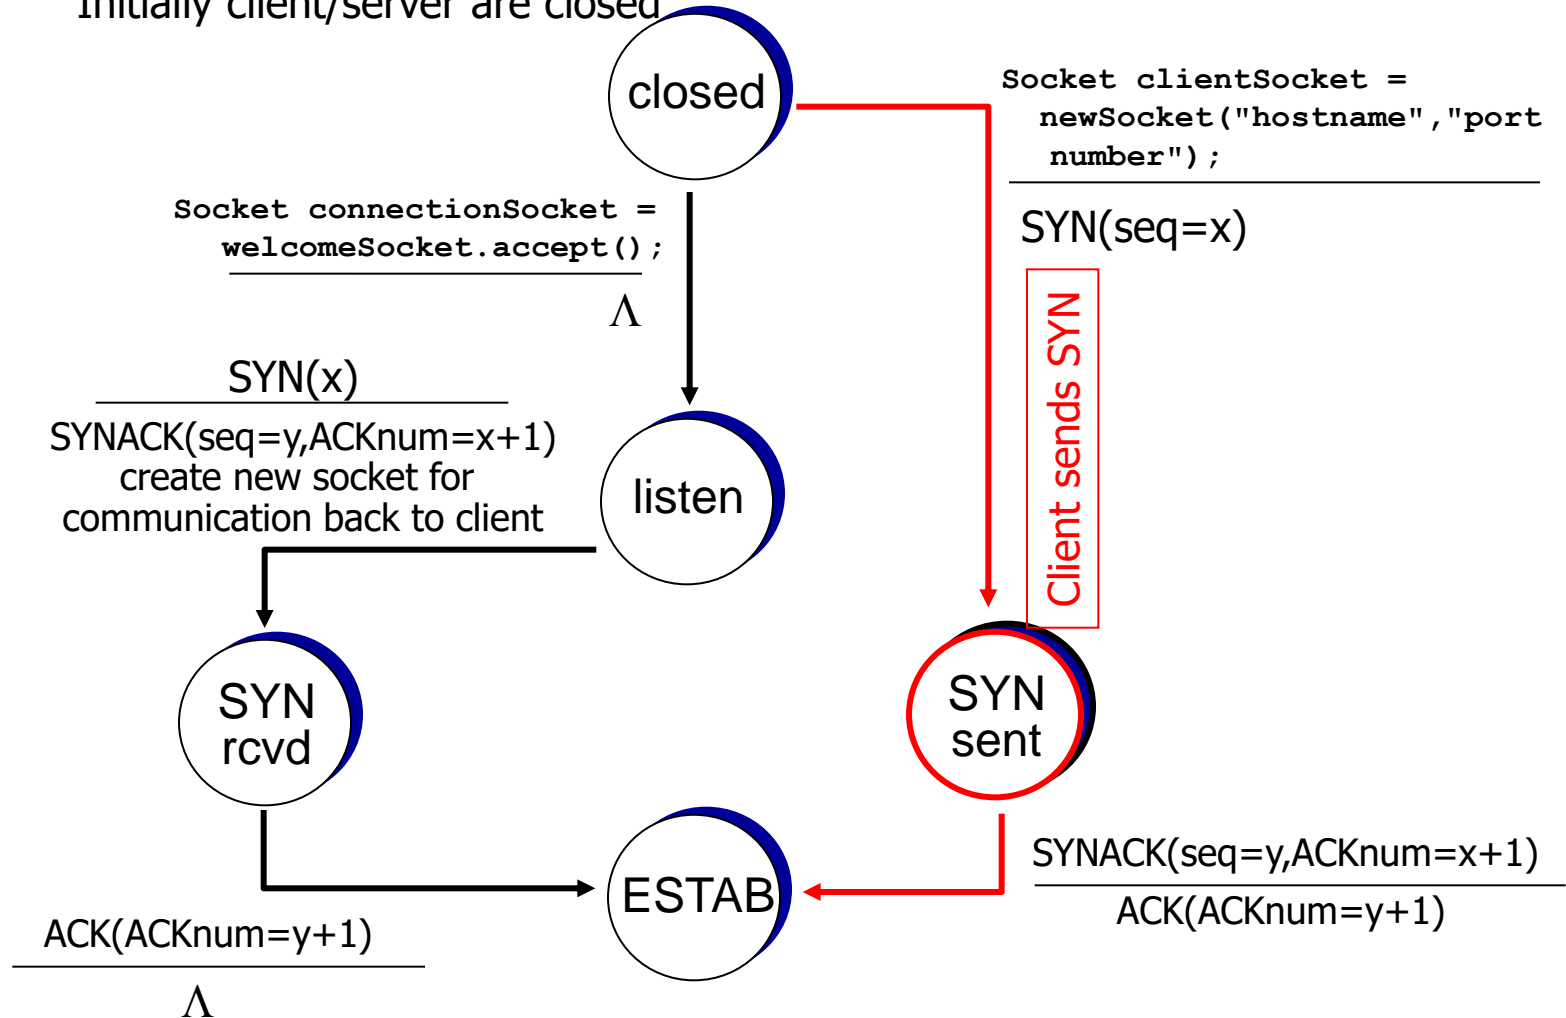


# TCP 3-way handshake



# TCP 3-way handshake: FSM

Initially client/server are closed



# TCP: closing a connection

- client, server each close their side of connection
  - send TCP segment with FIN bit = 1
- respond to received FIN with ACK
  - on receiving FIN, ACK can be combined with own FIN
- simultaneous FIN exchanges can be handled

# TCP: closing a connection

*client state*

ESTAB

`clientSocket.close()`

FIN\_WAIT\_1

can no longer  
send but can  
receive data

FIN\_WAIT\_2

wait for server  
close

TIMED\_WAIT

timed wait  
for  $2 * \text{max}$   
segment lifetime

CLOSED



FINbit=1, seq=x

ACKbit=1; ACKnum=x+1

FINbit=1, seq=y

ACKbit=1; ACKnum=y+1

can still  
send data

can no longer  
send data

*server state*

ESTAB

CLOSE\_WAIT

LAST\_ACK

CLOSED

# Chapter 3 outline

3.1 transport-layer services

3.2 multiplexing and demultiplexing

3.3 connectionless transport: UDP

3.4 principles of reliable data transfer

3.5 connection-oriented transport: TCP

- segment structure
- reliable data transfer
- flow control
- connection management

3.6 principles of congestion control

3.7 TCP congestion control